

CPSC 340: Machine Learning and Data Mining

Recommender Systems
Spring 2022 (2021W2)

Admin

- A5 due this Friday (midnight)
- Final exam format will be similar to the midterm. Online (Canvas).

Last Few Lectures: Latent-Factor Models

- We've been discussing latent-factor models of the form:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_i^T, z_j \rangle - x_{ij})^2$$

- We get different models under different conditions:
 - **K-means**: each z_i has one ‘1’ and the rest are zero.
 - **Least squares**: we only have one variable ($d=1$) and the z_i are fixed.
 - **PCA**: no restrictions on W or Z .
 - **Orthogonal PCA** (usual case): the rows w_c have norm 1 and inner products of zero.
 - **NMF**: all elements of W and Z are non-negative.

Variations on Latent-Factor Models


- We can use all our **tricks for linear regression** in this context:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d |w_j z_i - x_{ij}| + \frac{\lambda_1}{2} \sum_{i=1}^n \sum_{c=1}^k z_{ic}^2 + \frac{\lambda_2}{2} \sum_{j=1}^d \sum_{c=1}^k |w_{cj}|$$

- **Absolute loss** gives **robust PCA** that is less sensitive to outliers.
- We can use **L2-regularization**.
 - Though only reduces overfitting if we regularize both ‘W’ and ‘Z’.
- We can use **L1-regularization** to give sparse latent factors/features.
- We can use logistic/softmax/Poisson losses for discrete x_{ij} .
- Can use **change of basis** to learn **non-linear** latent-factor models.

bonus!



Application: Image Restoration



bonus!

Latent-Factor Models for Image Patches

- Consider building latent-factors for general **image patches**:




Size of X : $(\text{image height}) \times (\text{image width})$ by $((\text{patch height}) \times (\text{patch width}) \times 3)$

bonus!

Latent-Factor Models for Image Patches

- Consider building latent-factors for general **image patches**:




Typical pre-processing:

1. Usual variable centering
2. “Whiten” patches.
(remove correlations - bonus)

bonus!

Latent-Factor Models for Image Patches




(b) Principal components.

Orthogonal bases don't seem right:

- Few PCs do almost everything.
- Most PCs do almost nothing.

We believe “simple cells” in visual cortex use:




‘Gabor’ filters


bonus!

Latent-Factor Models for Image Patches

- Results from a “sparse” (non-orthogonal) latent factor model:



(a) With centering - gray.




(b) With centering - RGB.


bonus!

Latent-Factor Models for Image Patches

- Results from a “sparse” (non-orthogonal) latent-factor model:



(c) With whitening - gray.




(d) With whitening - RGB.

“colour
opponency”

bonus!

Recent Work: Structured Sparsity

- Basis learned with a variant of “structured sparsity”:




(b) With 4×4 neighborhood.

Similar to “cortical columns”
theory in visual cortex.

bonus!

Beyond NMF: Topic Models


- For modeling data as combinations of non-negative parts, NMF has largely been replaced by “topic models”.
 - A “fully-Bayesian” model where sparsity arises naturally.
 - Most popular example is called “latent Dirichlet allocation” (CPSC 440).



(pause)

Recommender System Motivation: Netflix Prize

- Netflix Prize:
 - 100M ratings from 0.5M users on 18k movies.
 - Grand prize was \$1M for first team to reduce squared error by 10%.
 - Started on October 2nd, 2006.
 - Netflix's system was first beat October 8th.
 - 1% error reduction achieved on October 15th.
 - Steady improvement after that.
 - ML methods soon dominated.
 - One obstacle was ‘Napoleon Dynamite’ problem:
 - Some movie ratings seem very difficult to predict.
 - Should only be recommended to certain groups.



Lessons Learned from Netflix Prize

- Prize awarded in 2009:
 - Ensemble method that averaged 107 models.
 - Increasing diversity of models more important than improving models.




- Winning entry (and most entries) used collaborative filtering:
 - Methods that only looks at ratings, not features of movies/users.
- A simple collaborative filtering method that does really well (7%):
 - “Regularized matrix factorization”. Now adopted by many companies.

Motivation: Other Recommender Systems

- Recommender systems are now everywhere:
 - Music, news, books, jokes, experts, restaurants, friends, dates, etc.
- Main types of approaches:
 1. Content-based filtering.
 - Supervised learning:
 - Extract features x_i of users and items, building model to predict rating y_i given x_i .
 - Apply model to prediction for new users/items.
 - Example: G-mail's "important messages" (personalization with "local" features).
 2. Collaborative filtering.
 - "Unsupervised" learning (have label matrix 'Y' but no features):
 - We only have labels y_{ij} (rating of user 'i' for movie 'j').
 - Example: Amazon recommendation algorithm.

Collaborative Filtering Problem


- Collaborative filtering is ‘filling in’ the **user-item matrix**:



- We have some ratings available with values {1, 2, 3, 4, 5}.
- We want to predict ratings “?” by looking at available ratings.

Collaborative Filtering Problem

- Collaborative filtering is ‘filling in’ the **user-item matrix**:



- What rating would “Ryan Reynolds” give to “Green Lantern”?
 - Why is this not completely hopeless? It *could* be anything.
 - But we may have **similar users and movies**.

Matrix Factorization for Collaborative Filtering

- Our standard **latent-factor model** for entries in matrix 'Y':

$$Y \approx ZW$$

$n \times d$ $n \times k$ $k \times d$

$$y_{ij} \approx \langle w^j, z_i \rangle$$

- User 'i' has latent features z_i .
- Movie 'j' has latent features w^j .
- Our loss functions sums over available ratings 'R':

$$f(Z, W) = \sum_{(i,j) \in R} (\langle w^j, z_i \rangle - y_{ij})^2 + \frac{\lambda_1}{2} \|Z\|_F^2 + \frac{\lambda_2}{2} \|W\|_F^2$$

- And we add **L2-regularization** to both types of features.
 - Basically, this is **regularized PCA** on the available entries of Y.
 - Typically fit with **SGD**.
- This simple method gives you a 7% improvement on the Netflix problem.

bonus!

Adding Global/User/Movie Biases

- Our standard **latent-factor model** for entries in matrix ‘Y’:

$$\hat{y}_{ij} = \langle w_j^j, z_i \rangle$$

- Sometimes we **don't assume the y_{ij} have a mean of zero**:

- We could add bias β reflecting average overall rating:

$$\hat{y}_{ij} = \beta + \langle w_j^j, z_i \rangle$$

- We could also add a **user-specific bias β_i** and **item-specific bias β_j** .

$$\hat{y}_{ij} = \beta + \beta_i + \beta_j + \langle w_j^j, z_i \rangle$$

- Some users are more generous, and some movies are just better.
 - These might also be regularized.

Beyond Accuracy in Recommender Systems

- Winning system of Netflix Challenge **was never adopted**.
- Other issues important in recommender systems:
 - **Diversity**: how different are the recommendations?
 - If you like ‘Battle of Five Armies Extended Edition’, recommend ‘Battle of Five Armies’?
 - Even if you really really like Star Wars, you might want non-Star-Wars suggestions.
 - **Persistence**: how long should recommendations last?
 - If you keep not clicking on ‘Justice League’, should it go away?
 - **Trust**: tell user *why* you made a recommendation.
 - **Social recommendation**: what did your friends watch?
 - **Freshness**: people tend to get more excited about *new/surprising* things.
 - Collaborative filtering does **not predict well for new users/movies**.
 - New movies don’t yet have ratings, and new users haven’t rated anything.



Content-Based vs. Collaborative Filtering

- Our latent-factor approach to **collaborative filtering** (Part 4):

$$\hat{y}_{ij} = \langle w^j, z_i \rangle$$

"hidden" features of movie w^j "hidden" features of user z_i

- Learns about each user/movie, but **can't predict on new users/movies**.
 - A linear model approach to **content-based filtering** (Part 3):
- $$\hat{y}_{ij} = w^T x_{ij}$$
- Here x_{ij} is a **vector of features** for the movie/user.
- Usual supervised learning setup: 'y' would contain all the y_{ij} , X would have x_{ij} as rows.
 - Can **predict on new users/movies**, but **can't learn about each user/movie**.
- Our usual supervised learning setup: $y_i = w^T x_i$

bonus!

Hybrid Approaches

- Hybrid approaches **combine content-based/collaborative filtering**:
 - SVDfeature (won “KDD Cup” in 2011 and 2012).

$$\hat{y}_{ij} = \beta + \beta_i + \beta_j + w^T x_{ij} + \langle w^i, z_i \rangle$$

Diagram annotations:

- Average rating across all users/movies (green bracket under β)
- Average rating for user ' i ' (green bracket under β_i)
- Average rating for movie ' j ' (green bracket under β_j)
- Linear model based on user/movie features x_{ij} . (blue bracket under $w^T x_{ij}$)
- Extra factors we learn for specific users and movies. (green bracket under $\langle w^i, z_i \rangle$)
- Latent features z_i for user ' i ' and latent features w^i for movie ' j '. (blue bracket under $\langle w^i, z_i \rangle$)
- Standard supervised learning: can predict for new users/movies (green bracket under the entire equation)

- Note that x_{ij} is a feature vector. Also, ‘ w ’ and ‘ w^i ’ are different parameters.

bonus!

Stochastic Gradient for SVDfeature

- Common approach to fitting SVDfeature is **stochastic gradient**.
- Previously you saw stochastic gradient for supervised learning:
 - Choose a random example ' i '
 - Update parameters ' w ' using gradient of example ' i '
- Stochastic gradient for SVDfeature (formulas as bonus):
 - Choose a random user ' i ' and a random product ' j '
 - Update β , β_i , β_j , w , z_i , and w^j based on their gradient for this user-product.

Updated every time

bonus!



Social Regularization

- Many recommenders are now connected to **social networks**.
 - “Login using your Facebook account”.
- Often, people like similar movies to their friends.
- Recent recommender systems use **social regularization**.
 - Add a “regularizer” encouraging friends’ weights to be similar:
$$\frac{\lambda}{2} \sum_{(i,j) \in \text{"friends"}} \|z_i - z_j\|^2$$
 - If we get a new user, recommendations are based on friend’s preferences.

(pause)


Latent-Factor Models for Visualization

- PCA takes features x_i and gives k -dimensional approximation z_i .
- If k is small, we can use this to visualize high-dimensional data.



Motivation for Non-Linear Latent-Factor Models

- But PCA is a **parametric linear** model
- PCA may not find obvious low-dimensional structure.



- We could use **change of basis** or **kernels**: but **still need to pick basis**.

Multi-Dimensional Scaling

- PCA for visualization:
 - We're using PCA to get the location of the z_i values.
 - We then plot the z_i values as locations in a scatterplot.
- Multi-dimensional scaling (MDS) is a crazy idea:
 - Let's directly optimize the pixel locations of the z_i values.
 - "Gradient descent on the points in a scatterplot".
 - Needs a "cost" function saying how "good" the z_i locations are.
 - Traditional MDS cost function:

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (||z_i - z_j|| - ||x_i - x_j||)^2$$

Try to make scatterplot distances match high-dimensional distance

Distance between points in original ' d ' dimensions



sum over pairs of examples

distance in scatterplot

Multi-Dimensional Scaling

- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.


$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (\|z_i - z_j\| - \|x_i - x_j\|)^2$$



t-Distributed Stochastic Neighbour Embedding

bonus!

- One key idea in t-SNE:
 - Focus on distance to “neighbours” (allow large variance in other distances)



Summary

- **Recommender systems** try to recommend products.
- **Collaborative filtering** tries to fill in missing values in a matrix.
 - Matrix factorization is a common approach.
- **Multi-dimensional scaling** is a non-parametric latent-factor model.
 - Big space of variants that we didn't have time to go into.
- Next time: the long-awaited start of deep learning.

bonus!

Digression: “Whitening”

- With image data, features will be very redundant.
 - Neighbouring pixels tend to have similar values.
- A standard transformation in these settings is “**whitening**”:
 - Rotate the data so features are uncorrelated.
 - Re-scale the rotated features so they have a variance of 1.
- Using SVD approach to PCA, we can do this with:
 - Get ‘W’ from SVD (usually with $k=d$).
 - $Z = XW^T$ (rotate to give uncorrelated features).
 - Divide columns of ‘Z’ by corresponding singular values (unit variance).
- Details/discussion [here](#).

bonus!

Motivation for Topic Models

- Want a model of the “factors” making up documents.
 - Instead of latent-factor models, they’re called **topic models**.
 - The canonical topic model is **latent Dirichlet allocation (LDA)**.

Suppose you have the following set of sentences:

- I like to eat broccoli and bananas.
- I ate a banana and spinach smoothie for breakfast.
- Chinchillas and kittens are cute.
- My sister adopted a kitten yesterday.
- Look at this cute hamster munching on a piece of broccoli.

What is latent Dirichlet allocation? It’s a way of automatically discovering **topics** that these sentences contain. For example, given these sentences and asked for 2 topics, LDA might produce something like

- **Sentences 1 and 2:** 100% Topic A
- **Sentences 3 and 4:** 100% Topic B
- **Sentence 5:** 60% Topic A, 40% Topic B
- **Topic A:** 30% broccoli, 15% bananas, 10% breakfast, 10% munching, ... (at which point, you could interpret topic A to be about food)
- **Topic B:** 20% chinchillas, 20% kittens, 20% cute, 15% hamster, ... (at which point, you could interpret topic B to be about cute animals)

- “Topics” could be useful for things like searching for relevant documents.

bonus!

Term Frequency – Inverse Document Frequency

- In information retrieval, classic word importance measure is **TF-IDF**.
- First part is the **term frequency** $tf(t,d)$ of term ‘t’ for document ‘d’.
 - Number of times “word” ‘t’ occurs in document ‘d’, divided by total words.
 - E.g., 7% of words in document ‘d’ are “the” and 2% of the words are “Lebron”.
- Second part is **document frequency** $df(t,D)$.
 - Compute number of documents that have ‘t’ at least once.
 - E.g., 100% of documents contain “the” and 0.01% have “LeBron”.
- TF-IDF is $tf(t,d) * \log(1/df(t,D))$.

bonus!

Term Frequency – Inverse Document Frequency

- The **TF-IDF** statistic is $tf(t,d) * \log(1/df(t,D))$.
 - It's high if word 't' happens often in document 'd', but isn't common.
 - E.g., seeing "LeBron" a lot it tells you something about "topic" of article.
 - E.g., seeing "the" a lot tells you nothing.
- There are ***many*** variations on this statistic.
 - E.g., avoiding dividing by zero and all types of "frequencies".
- Summarizing 'n' documents into a matrix X:
 - Each row corresponds to a document.
 - Each column gives the TF-IDF value of a particular word in the document.

bonus!

Latent Semantic Indexing

- TF-IDF features are **very redundant**.
 - Consider TF-IDFs of “LeBron”, “Durant”, “Harden”, and “Kobe”.
 - High values of these typically just indicate topic of “basketball”.
- We can probably compress this information quite a bit.
- Latent Semantic Indexing/Analysis:
 - Run **latent-factor model** (like PCA or NMF) on TF-IDF matrix X.
 - Treat the principal components as the “topics”.
 - **Latent Dirichlet allocation** is a variant that avoids weird $df(t,D)$ heuristic.

bonus!

SVDfeature with SGD: the gory details

Objective: $\frac{1}{2} \sum_{(i,j) \in R} (\hat{y}_{ij} - y_{ij})^2$ with $\hat{y}_{ij} = \beta + \beta_i + \beta_j + w^T x_{ij} + (w^j)^T z_i$

Update based on random (i, j) :

$$\beta = \beta - \alpha r_{ij}$$

$$\beta_i = \beta_i - \alpha r_{ij}$$

$$\beta_j = \beta_j - \alpha r_{ij}$$

but, updates are the same,
 'beta' is always update while β_i and β_j are
 only updated for the specific user + product

$$w = w - \alpha r_{ij} x_{ij} \leftarrow \text{updated every time.}$$

$$z_i = z_i - \alpha r_{ij} w^j$$

$$w^j = w^j - \alpha r_{ij} z_i$$

updated for
 specific user
 and product.

(Adding regularization adds an extra term)

bonus!

Tensor Factorization


- Tensors are higher-order generalizations of matrices:

Scalar $\alpha = []_{1 \times 1}$

Vector $a = []_{d \times 1}$

Matrix $A = []_{d \times d}$

Tensor $A = []_{d \times d \times d}$



- Generalization of matrix factorization is **tensor factorization**:

$$y_{ijm} \approx \sum_{c=1}^k w_{jc} z_{ic} v_{mc}$$

- Useful if there are other relevant variables:

- Instead of ratings based on {user,movie}, ratings based {user,movie,group}.
- Useful if you have groups of users, or if ratings change over time.

bonus!

Field-Aware Matrix Factorization

- Field-aware factorization machines (FFMs):
 - Matrix factorization with multiple z_i or w_c for each example or part.
 - You choose which z_i or w_c to use based on the value of feature.
- Example from “click through rate” prediction:
 - E.g., predict whether “male” clicks on “nike” advertising on “espn” page.
 - A previous matrix factorization method for the 3 factors used:
$$w_{espn} w_{nike} + w_{espn} w_{male} + w_{nike} w_{male}$$
$$w_{espn}^A w_{nike}^P + w_{espn}^G w_{male}^P + w_{nike}^G w_{male}^A$$
 - FFM could use:
 - $w_{espn}A$ is the factor we use when multiplying by a an advertiser’s latent factor.
 - $w_{espn}G$ is the factor we use when multiplying by a group’s latent factor.
- This approach has won some Kaggle competitions ([link](#)), and has shown to work well in production systems too ([link](#)).

bonus!

Warm-Starting

- We've used data $\{X, y\}$ to fit a model.
- We now have new training data and want to fit new and old data.
- Do we need to re-fit from scratch?
- This is the warm starting problem.
 - It's easier to warm start some models than others.

Easy Case: K-Nearest Neighbours and Counting

bonus!

- **K-nearest neighbours:**

- KNN just stores the training data, so just store the new data.

- **Counting-based** models:

- Models that base predictions on frequencies of events.
 - E.g., naïve Bayes.

- Just update the counts:

$$p(\text{"vicodin"} \mid \text{"spam"}) = \frac{\text{count of } \{\text{"vicodin"}, \text{"spam"}\} \text{ in } \underline{\text{new and old data}}}{\text{count of } \text{"spam"} \text{ in } \underline{\text{new and old data}}}$$

- Decision trees with fixed rules: just update counts at the leaves.

bonus!

Medium Case: L2-Regularized Least Squares

- L2-regularized least squares is obtained from linear algebra:

$$w = (X^T X + \lambda I)^{-1} (X^T y)$$

- Cost is $O(nd^2 + d^3)$ for ‘n’ training examples and ‘d’ features.
- Given one new point, we need to compute:
 - $X^T y$ with one row added, which costs $O(d)$.
 - Old $X^T X$ plus $x_i x_i^T$, which costs $O(d^2)$.
 - Solution of linear system, which costs $O(d^3)$.
 - So cost of adding ‘t’ new data point is $O(td^3)$.
- With “matrix factorization updates”, can reduce this to $O(td^2)$.
 - Cheaper than computing from scratch, particularly for large d.


bonus!

Medium Case: Logistic Regression

- We fit logistic regression by gradient descent on a convex function.
- With new data, convex function $f(w)$ changes to new function $g(w)$.

$$f(w) = \sum_{i=1}^n f_i(w) \quad g(w) = \sum_{i=1}^{n+1} f_i(w)$$

- If we don't have much more data, 'f' and 'g' will be "close".
 - Start gradient descent on 'g' with minimizer of 'f'.
 - You can show that it requires fewer iterations.



bonus!

Hard Cases: Non-Convex/Greedy Models

- For decision trees:
 - “Warm start”: continue splitting nodes that haven’t already been split.
 - “Cold start”: re-fit everything.
- Unlike previous cases, this **won’t in general give same result as re-fitting**:
 - New data points might lead to **different splits** higher up in the tree.
- Intermediate: usually do warm start but occasionally do a cold start.
- Similar heuristics/conclusions for other non-convex/greedy models:
 - K-means clustering.
 - Matrix factorization (though you can continue PCA algorithms).