

CPSC 340: Machine Learning and Data Mining

Deep Learning
Spring 2022 (2021W2)

Admin

- Course surveys
 - Please fill them out
 - We care deeply about your education, so we take them very seriously
 - You will be able to evaluate the class overall, and then Mijung and I separately
 - Please use the text boxes to also let us know about the “lecture specialization experiment” [where we each specialized in half the lectures]
 - As always, please remember we’re real people, so both praise and critical feedback are great. Please avoid personal, hurtful, or unconstructive negative comments.
- A5 deadline tonight
- **A6** out soon: by Monday at the latest, due April 8 (our last class)

End of Part 4: Key Concepts

- We discussed **linear latent-factor models**:

$$\begin{aligned} f(W, z) &= \sum_{i=1}^n \sum_{j=1}^d (\langle w_j^T z_i \rangle - x_{ij})^2 \\ &= \sum_{i=1}^n \|W^T z_i - x_i\|^2 \\ &= \|Z^T W - X\|_F^2 \end{aligned}$$

- Represent 'X' as linear combination of **latent factors 'w_c'**.
 - **Latent features 'z_i'** give a lower-dimensional version of each 'x_i'.
 - When k=1, finds **direction that minimizes squared orthogonal distance**.
- Applications:
 - Outlier detection, dimensionality reduction, data compression, features for linear models, visualization, factor discovery, filling in missing entries.

End of Part 4: Key Concepts

- We discussed linear latent-factor models:

$$f(W, z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_j^T z_i \rangle - x_{ij})^2$$

- Principal component analysis (PCA):

- Often uses orthogonal factors and fits them sequentially (via SVD).

- Non-negative matrix factorization:

- Uses non-negative factors giving sparsity.
 - Can be minimized with projected gradient.

- Many variations are possible:

- Different regularizers (sparse coding) or loss functions (robust/binary PCA).
 - Missing values (recommender systems) or change of basis (kernel PCA).

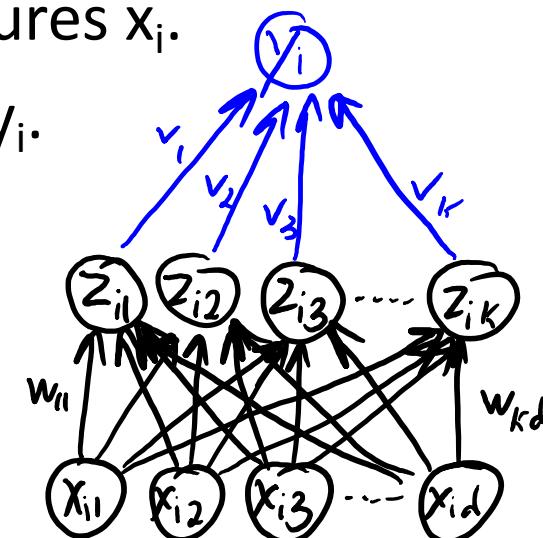
bonus!

End of Part 4: Key Concepts

- We didn't really discuss **multi-dimensional scaling (MDS)**:
 - Non-parametric method for high-dimensional data visualization.
 - Tries to match distance/similarity in high-/low-dimensions.
 - “Gradient descent on scatterplot points”.
- Main challenge in MDS methods is “crowding” effect:
 - Methods focus on large distances and lose local structure.
- Common solutions:
 - Sammon mapping: use weighted cost function.
 - ISOMAP: approximate geodesic distance using via shortest paths in graph.
 - T-SNE: give up on large distances and focus on neighbour distances.
- Word2vec is a recent MDS method giving better “word features”.

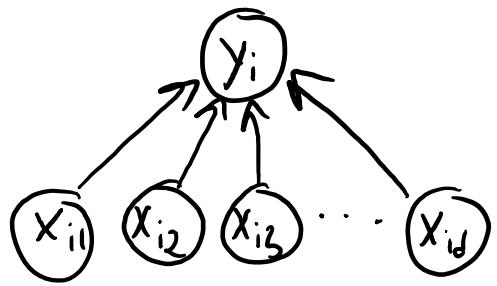
Supervised Learning Roadmap

- Part 1: “Direct” Supervised Learning.
 - We learned parameters ‘w’ based on the original features x_i and target y_i .
- Part 3: Change of Basis.
 - We learned parameters ‘v’ based on a change of basis z_i and target y_i .
- Part 4: Latent-Factor Models.
 - We learned parameters ‘W’ for basis z_i based on only on features x_i .
 - You can then learn ‘v’ based on change of basis z_i and target y_i .
- Part 5: Neural Networks.
 - Jointly learn ‘W’ and ‘v’ based on x_i and y_i .
 - Learn basis z_i that is good for supervised learning.

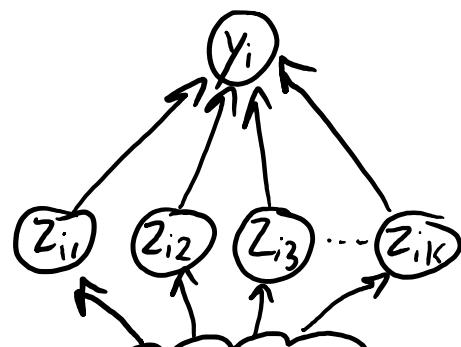


A Graphical Summary of CPSC 340 Parts 1-5

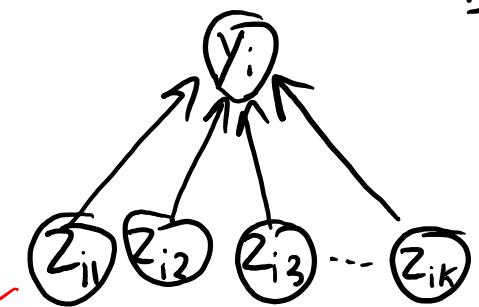
Part 1: "I have features x_i "



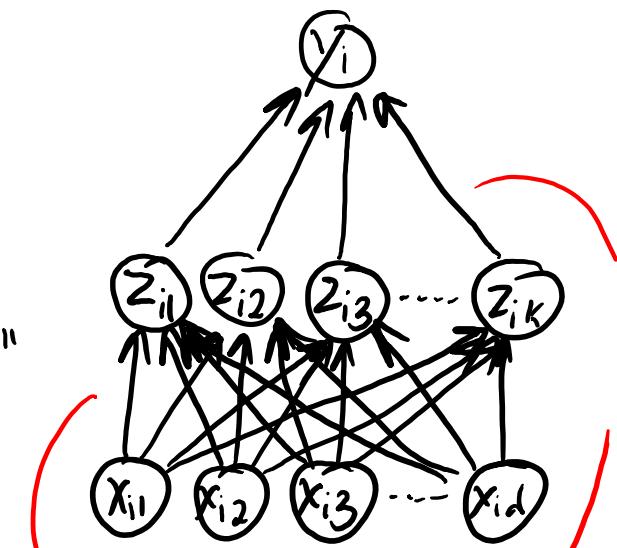
Part 3: Change of basis



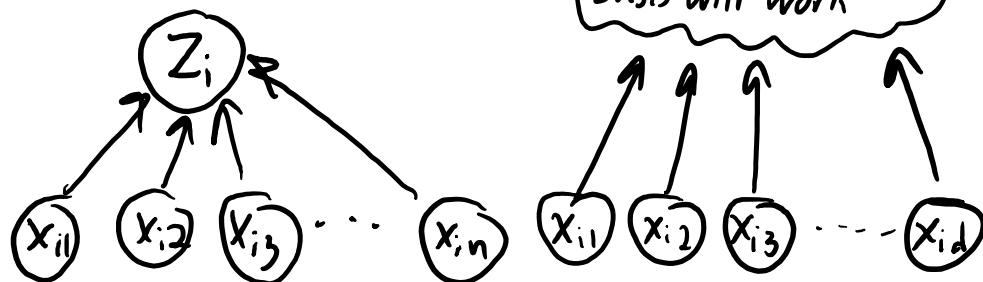
Part 4: basis from latent-factor model



Part 5: Neural networks



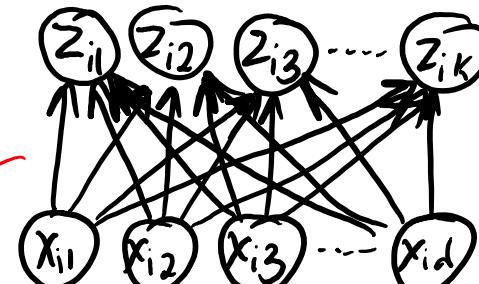
Part 2: "What is the group of x_i ?"



Trained separately

"What are the 'parts' of x_i ?"

"PCA will give me good features"



Learn features and classifier at the same time.

Notation for Neural Networks (MEMORIZE)

We have our usual supervised learning notation:

$$X = \begin{bmatrix} \cdots & x_1^T & \cdots \\ \cdots & x_2^T & \cdots \\ \vdots & & \vdots \\ \cdots & x_n^T & \cdots \end{bmatrix} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$n \times d$ $n \times 1$

We have our latent features:

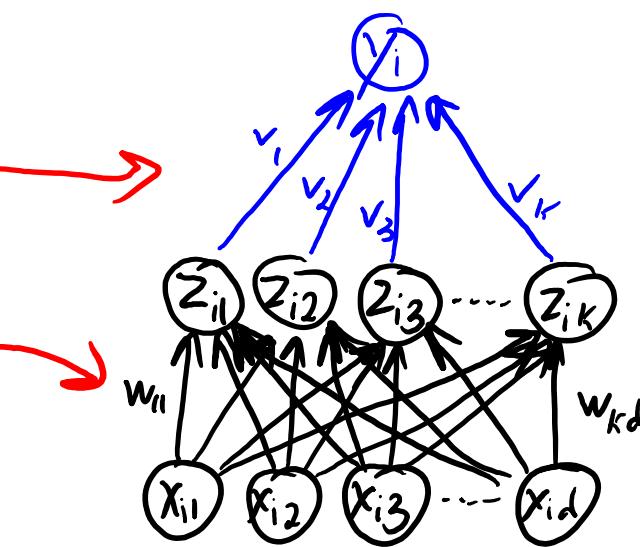
$$Z = \begin{bmatrix} \cdots & z_1^T & \cdots \\ \cdots & z_2^T & \cdots \\ \vdots & & \vdots \\ \cdots & z_n^T & \cdots \end{bmatrix}$$

$n \times K$

We have two sets of parameters:

$$V = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_K \end{bmatrix} \quad W = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_K \end{bmatrix}$$

$K \times 1$ $K \times d$



Linear-Linear Model

- Natural choice: linear latent-factor model with linear regression.

Use features from latent-factor model: $z_i = Wx_i$

Make predictions using a linear model: $y_i = v^T z_i$

- We want to train 'W' and 'v' jointly, so we could minimize:

$$f(W, v) = \frac{1}{2} \sum_{i=1}^n (v^T z_i - y_i)^2 = \frac{1}{2} \sum_{i=1}^n (v^T (Wx_i) - y_i)^2$$

linear regression
with z_i as features z_i come from
latent-factor model

- But this is just a linear model:

$$y_i = v^T z_i = v^T (Wx_i) = (\underbrace{v^T W}_{\text{some vector } 'w'}) x_i = w^T x_i$$

$1 \times d$

Introducing Non-Linearity

- To increase flexibility, something needs to be non-linear.
- Typical choice: transform z_i by non-linear function ‘ h ’.

$$z_i = Wx_i \quad \hat{y}_i = v^T h(z_i)$$

- Here the function ‘ h ’ transforms ‘ k ’ inputs to ‘ k ’ outputs.
- Common choice for ‘ h ’: applying sigmoid function element-wise:

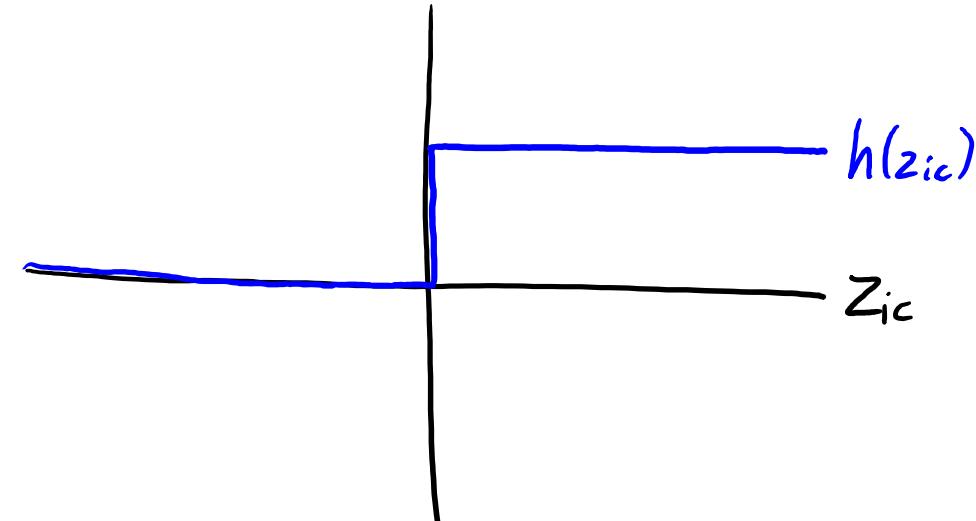
$$h(z_{ic}) = \frac{1}{1 + \exp(-z_{ic})}$$

- So this takes the z_{ic} in $(-\infty, \infty)$ and maps it to $(0, 1)$.
- This is called a “multi-layer perceptron” or a “neural network”.

Why Sigmoid?

- Consider setting ‘ h ’ to define **binary features** z_i using:

$$h(z_{ic}) = \begin{cases} 1 & \text{if } z_{ic} \geq 0 \\ 0 & \text{if } z_{ic} < 0 \end{cases}$$



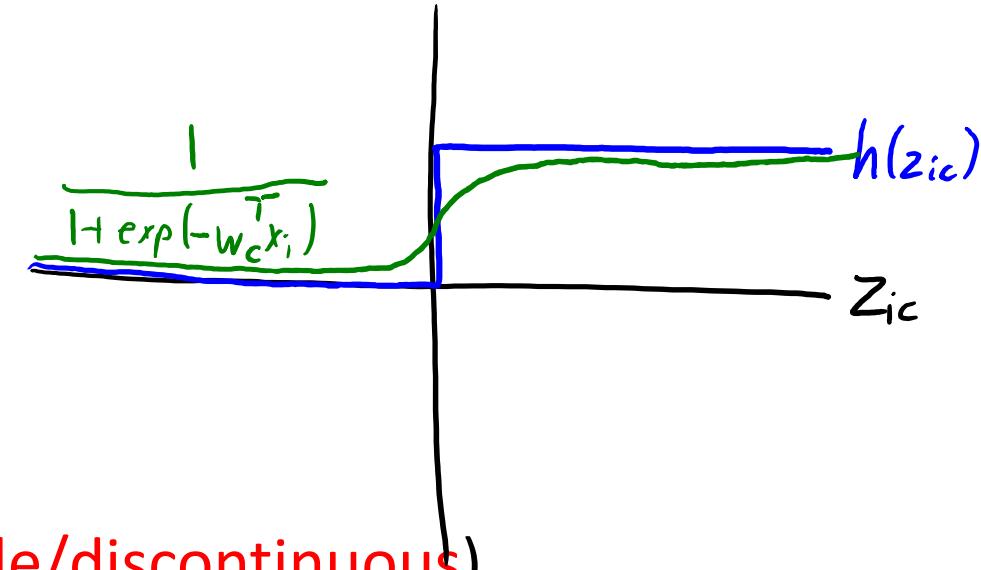
- Each $h(z_i)$ can be viewed as binary feature.
 - “You either have this ‘part’ or you don’t have it.”
- We can make 2^k objects by all the possible “part combinations”.

| Motivation: Pixels vs. Parts | | | | | | | | | | |
|---|-----|---|-----|---|-----|---|-----|---|-----|---|
| • We could represent other digits as different combinations of “parts”: | | | | | | | | | | |
| 3 | = 1 | - | + 1 | 1 | + 1 | - | + 1 | 1 | + 1 | 0 |
| 5 | = 1 | - | + 0 | 1 | + 1 | - | + 1 | 1 | + 0 | 1 |
| 8 | = 1 | - | + 1 | 1 | + 1 | - | + 1 | 1 | + 1 | 1 |

Why Sigmoid?

- Consider setting ‘ h ’ to define **binary features z_i** using:

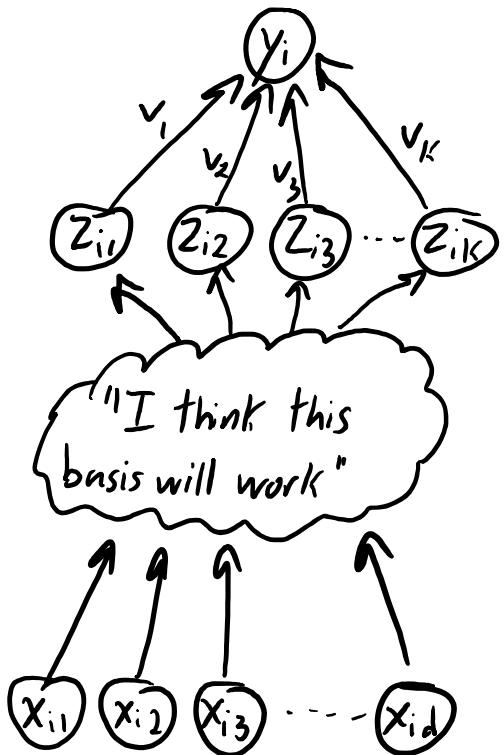
$$h(z_{ic}) = \begin{cases} 1 & \text{if } z_{ic} \geq 0 \\ 0 & \text{if } z_{ic} < 0 \end{cases}$$



- Each $h(z_i)$ can be viewed as binary feature.
 - “You either have this ‘part’ or you don’t have it.”
- But this is hard to optimize (**non-differentiable/discontinuous**).
- **Sigmoid is a smooth approximation** to these binary features.
 - Non-parametric version is a **universal approximator**:
 - If ‘ k ’ grows appropriately with ‘ n ’, can model any continuous function.

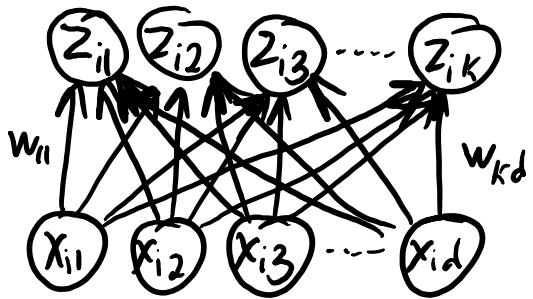
Supervised Learning Roadmap

Hand-engineered features:

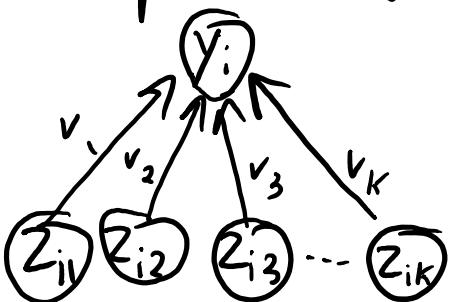


Requires domain knowledge
and can be time-consuming

Learn a latent-factor model:

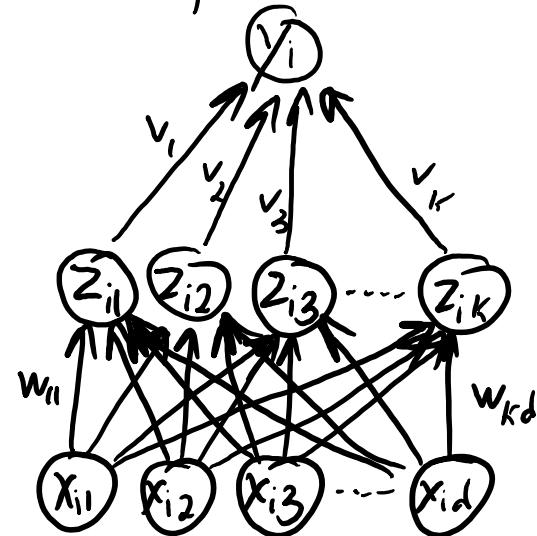


Use latent features
in supervised model:



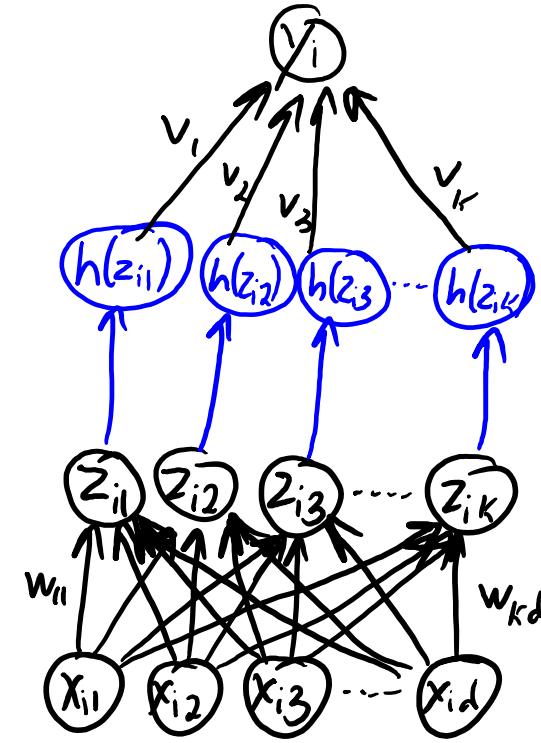
Good representation of
 x_i might be bad for predicting y_i

Learn ' w ' and ' W '
together:



But still gives a
linear model.

Neural network:

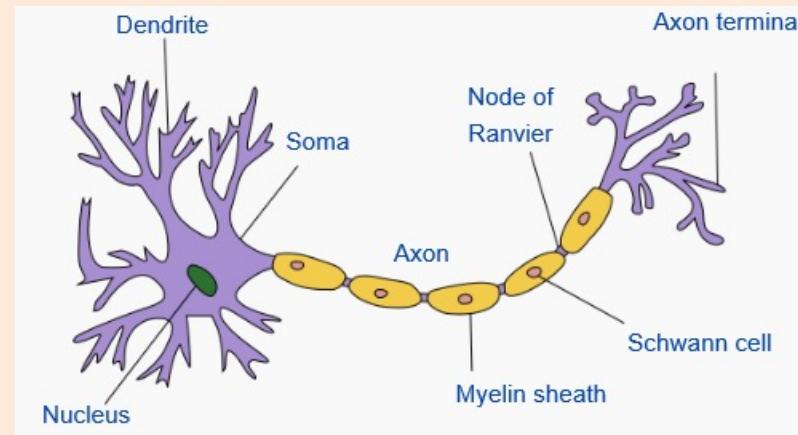


Extra non-linear
transformation ' h '!

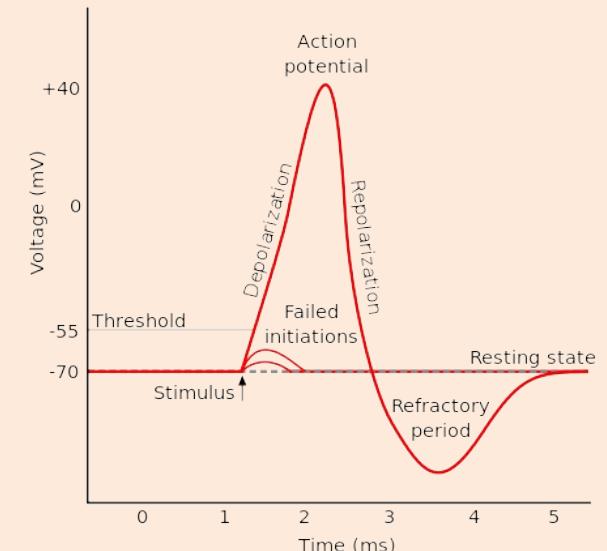
bonus!

Why “Neural Network”?

- Cartoon of “typical” neuron:

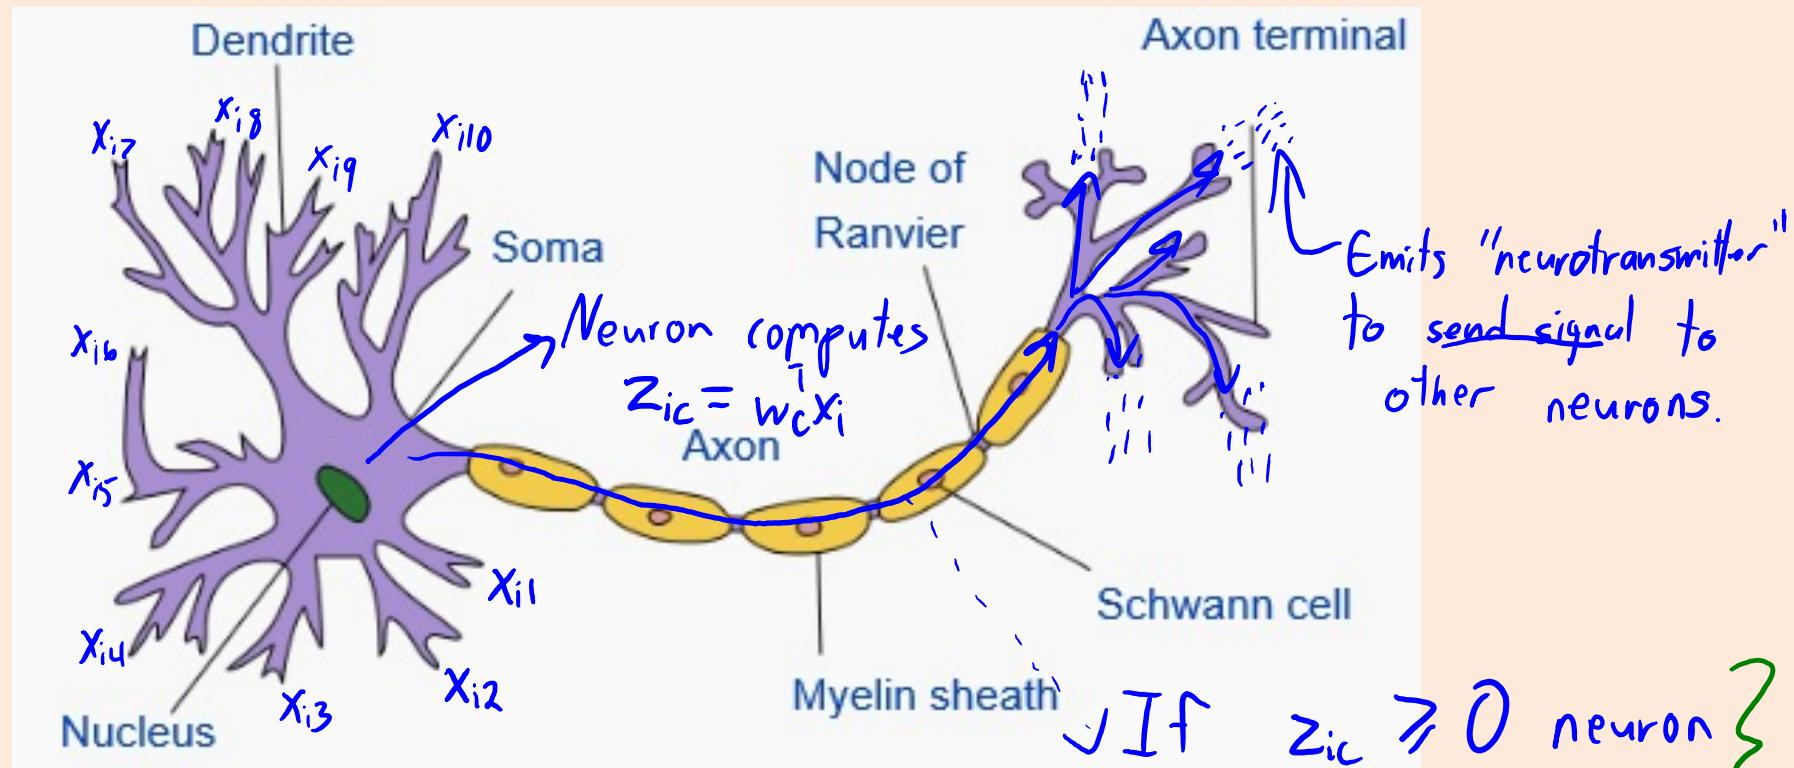


- Neuron has many “dendrites”, which take an input signal.
- Neuron has a single “axon”, which sends an output signal.
- With the right input to dendrites:
 - “Action potential” along axon (like a binary signal):

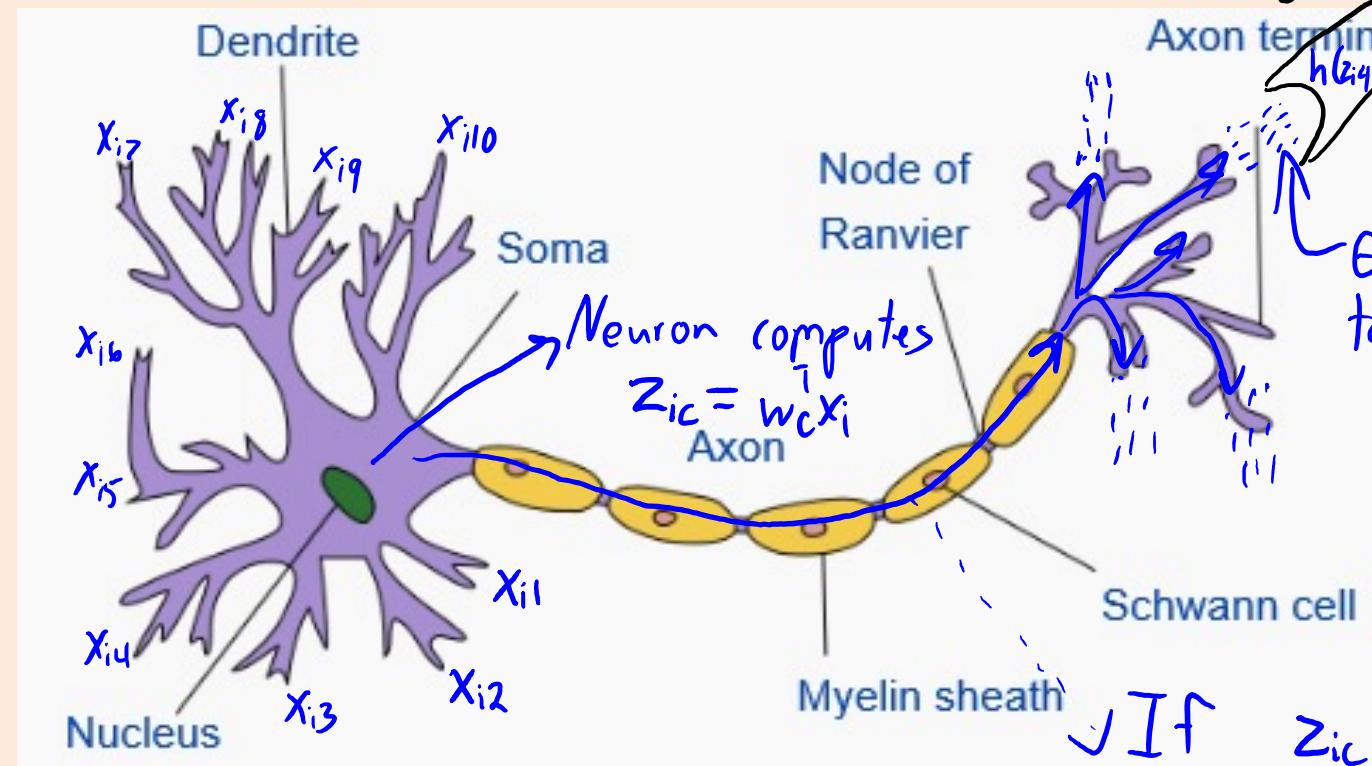


bonus!

Why “Neural Network”?

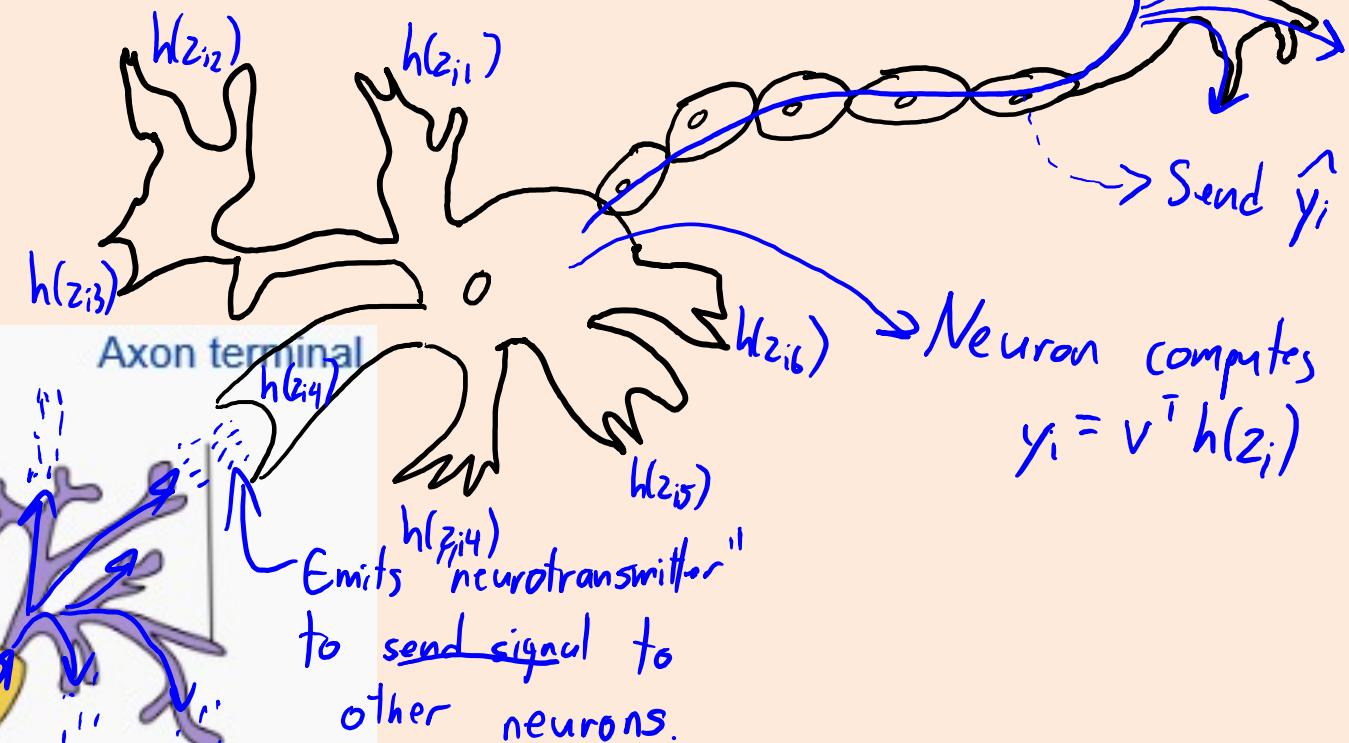


Why “Neural Network”?



If $Z_{ic} \geq 0$ neuron
Sends signal along axon.

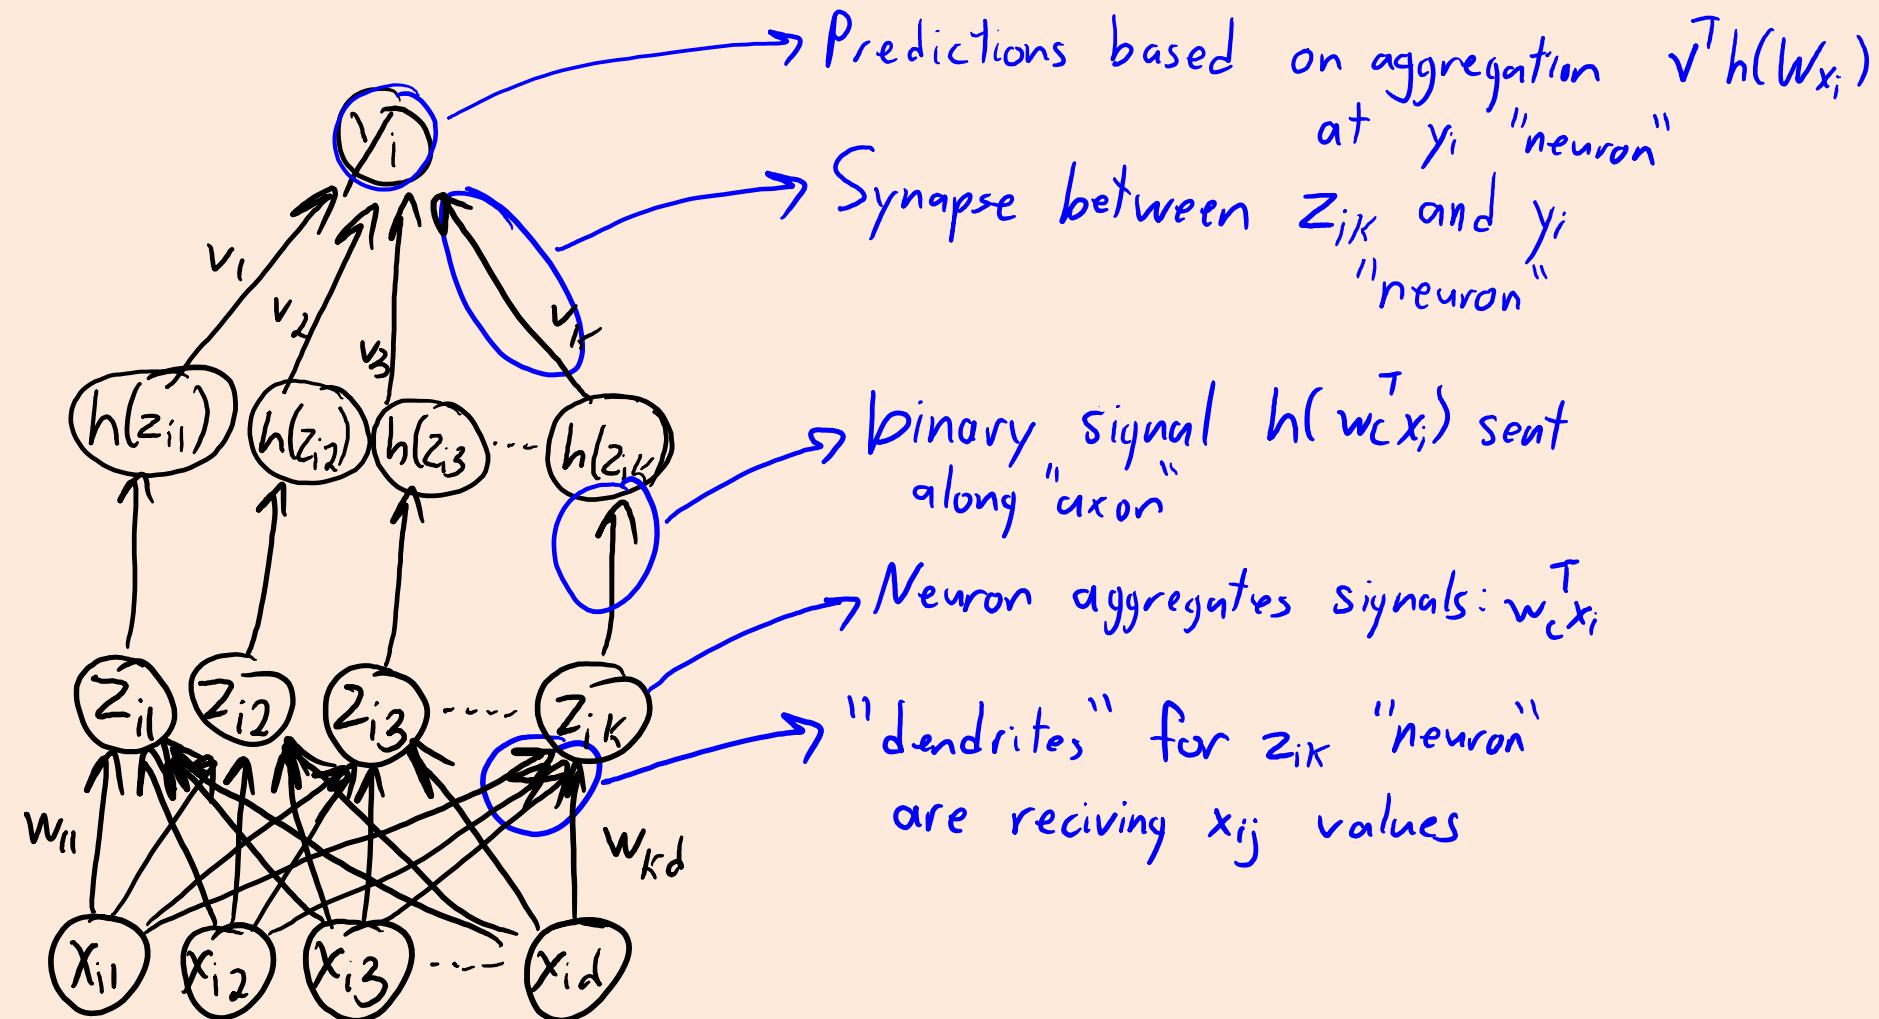
We approximate binary
signal with $\frac{1}{1 + \exp(-Z_{ic})}$



bonus!

bonus!

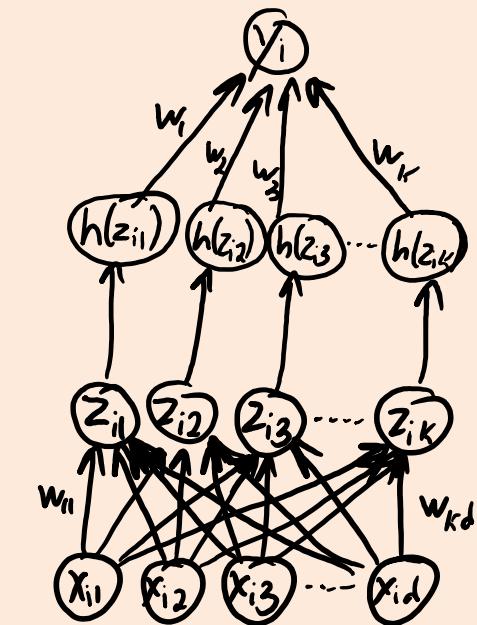
Why “Neural Network”?



bonus!

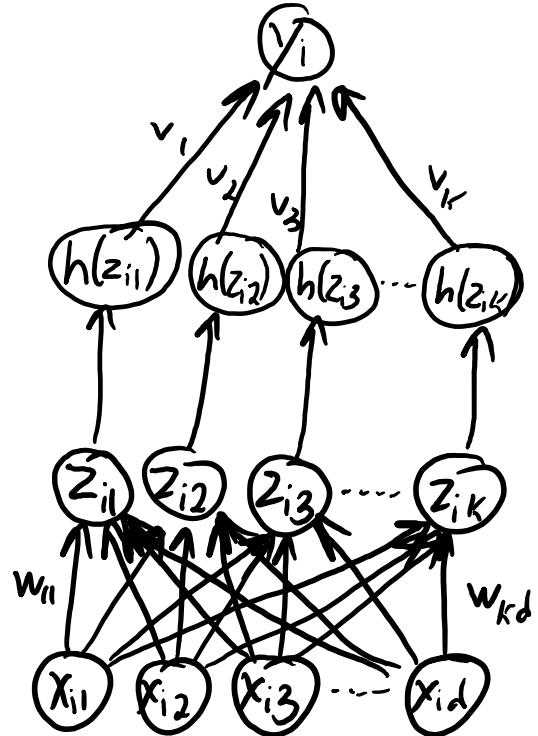
“Artificial” Neural Nets vs. “Real” Networks Nets

- Artificial neural network:
 - x_i is measurement of the world.
 - z_i is internal representation of world.
 - y_i is output of neuron for classification/regression.
- Real neural networks are more complicated:
 - **Timing** of action potentials seems to be important.
 - “Rate coding”: frequency of action potentials simulates continuous output.
 - Neural networks don’t reflect **sparsity** of action potentials.
 - How much computation is done **inside neuron?**
 - Brain is highly **organized** (e.g., substructures and cortical columns).
 - Connection **structure changes**.
 - **Different types** of neurotransmitters.



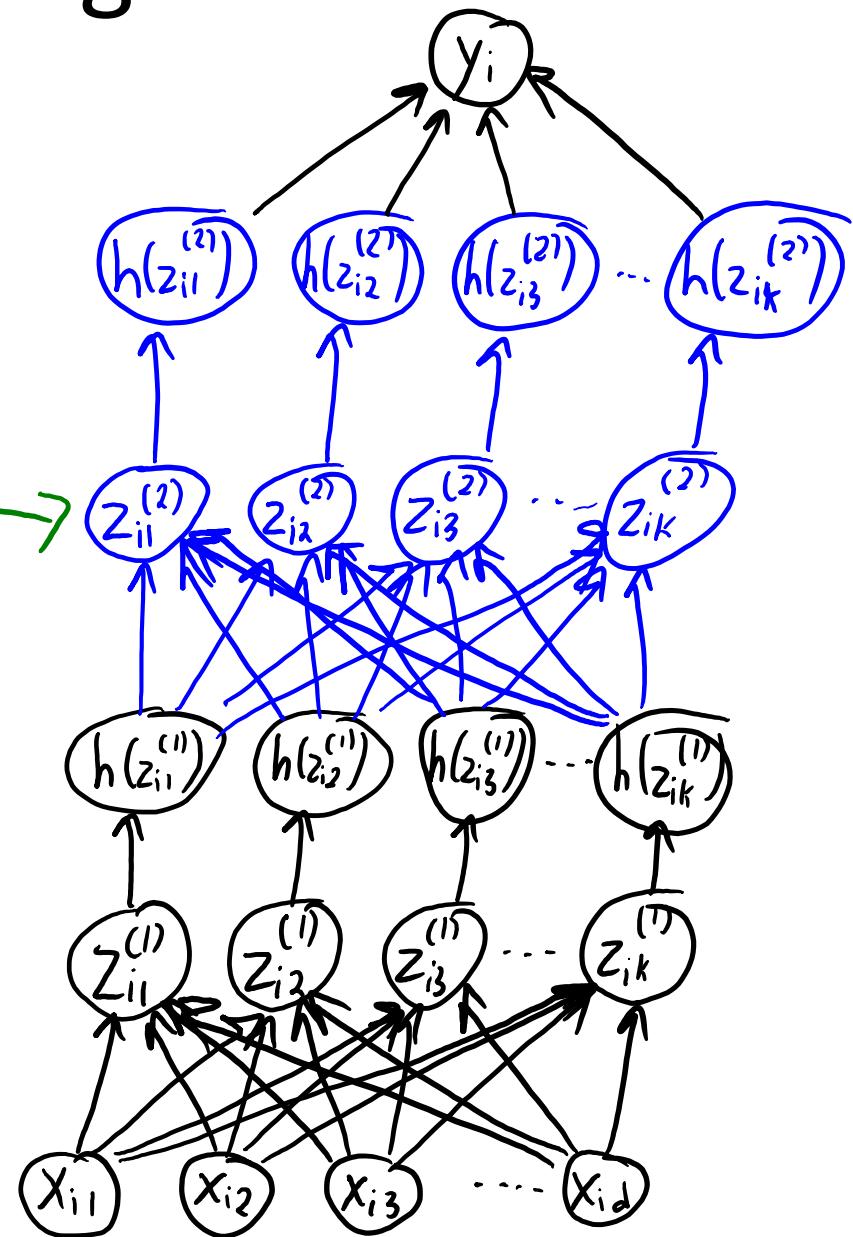
Deep Learning

Neural network:



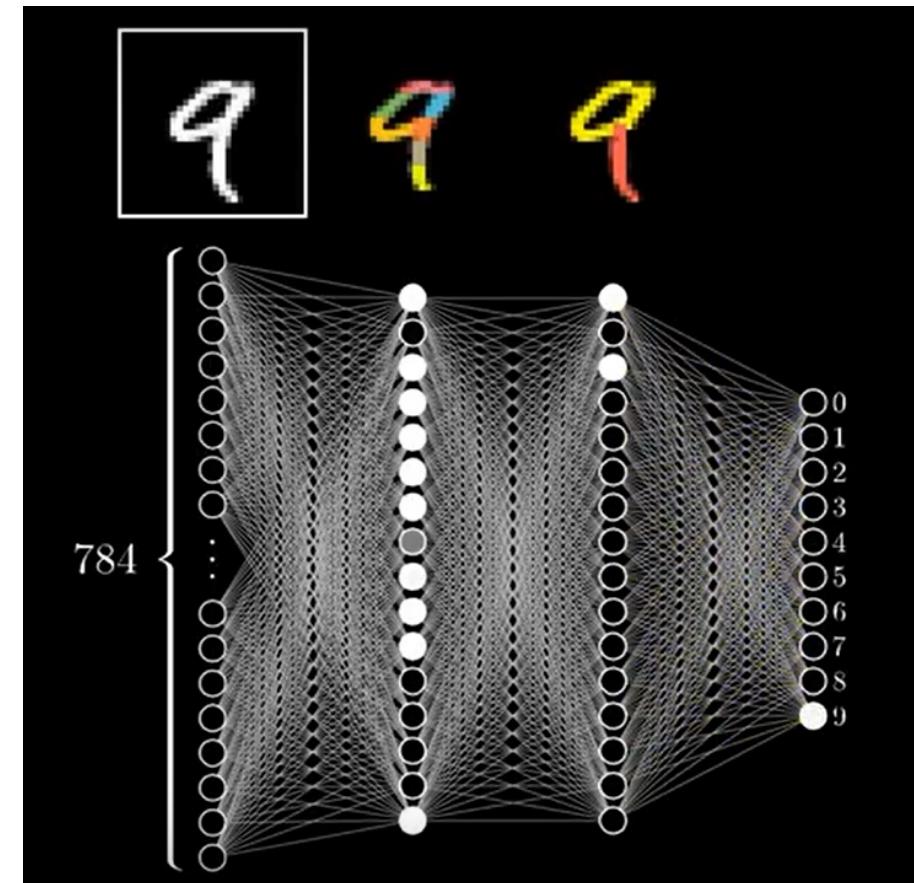
Deep learning:

Second "layer" of latent features
↑
You can add more "layers" to go "deeper"



“Hierarchies of Parts” Motivation for Deep Learning

- Each “neuron” might recognize a “part” of a digit.
 - “Deeper” neurons might recognize combinations of parts.
 - Represent complex objects as hierarchical combinations of re-useable parts (a simple “grammar”).
- Watch the full video here:
 - <https://www.youtube.com/watch?v=aircAruvnKk>
- Theory:
 - 1 big-enough hidden layer already gives universal approximation.
 - But some functions require exponentially-fewer parameters to approximate with more layers (can fight curse of dimensionality).



Deep Learning

Linear model:

$$\hat{y}_i = w^\top x_i$$

Neural network with 1 hidden layer:

$$\hat{y}_i = v^\top h(Wx_i)$$

Neural network with 2 hidden layers:

$$\hat{y}_i = v^\top h(W^{(2)} h(W^{(1)} x_i))$$

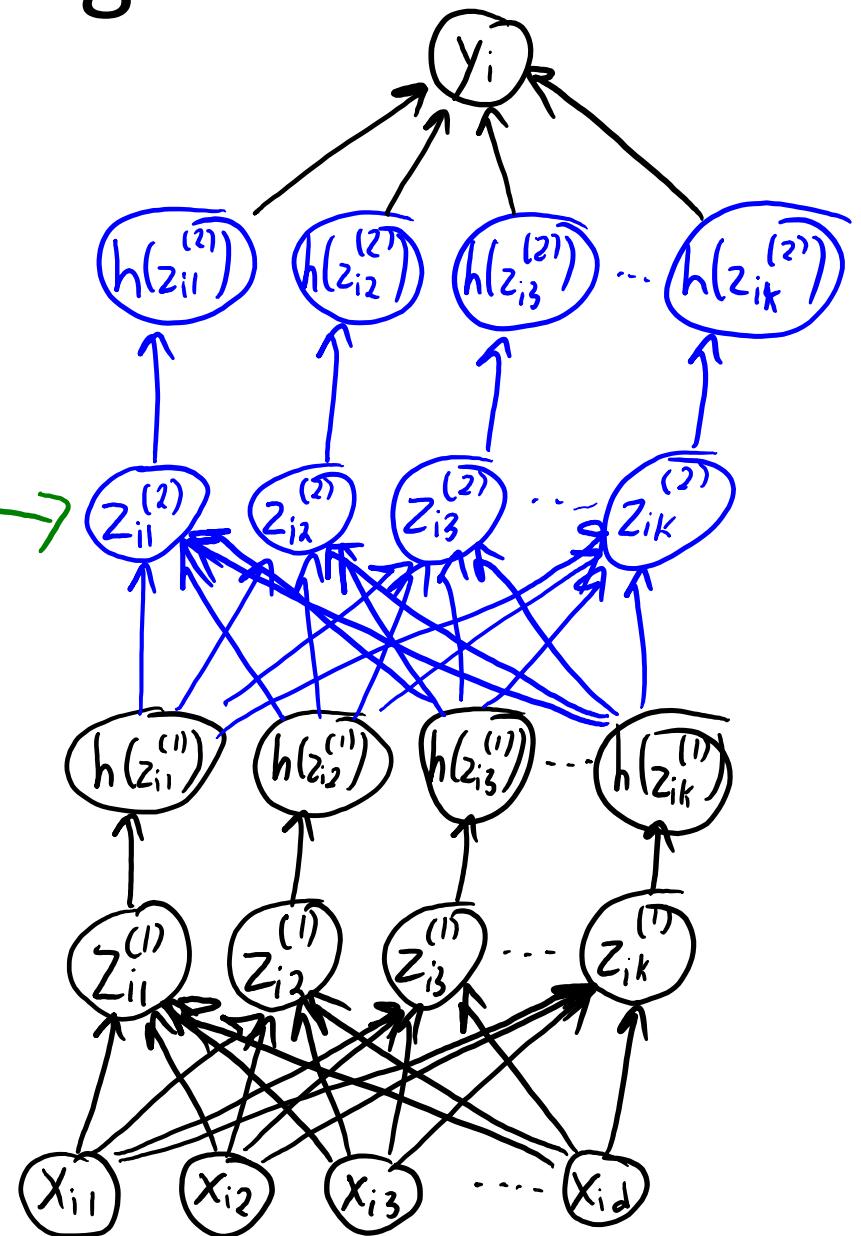
Neural network with 3 hidden layers

$$\hat{y}_i = v^\top h(W^{(3)} h(W^{(2)} h(W^{(1)} x_i)))$$

Deep learning:

Second "layer" of latent features

You can add more "layers" to go "deeper"



Deep Learning

- For 4 layers, we could write the prediction as:

$$\hat{y}_i = v^\top h(w^{(4)} h(w^{(3)} h(w^{(2)} h(w^{(1)} x_i))))$$

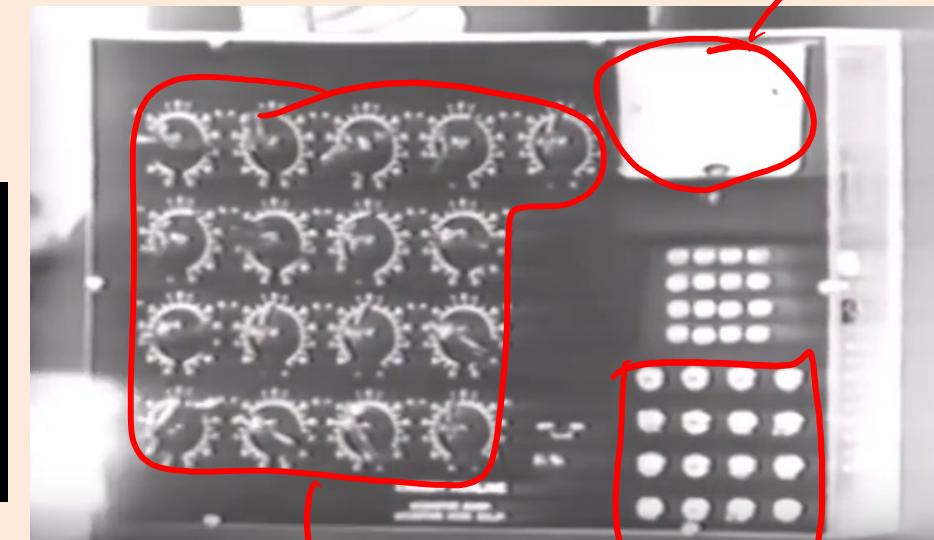
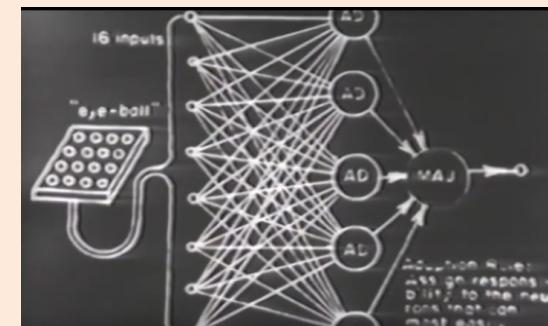
- For 'm' layers, we usually just say:

$$\hat{y}_i = v^\top h(w^{(m)} h(\dots h(w^{(1)} x_i)))$$

bonus!

ML and Deep Learning History

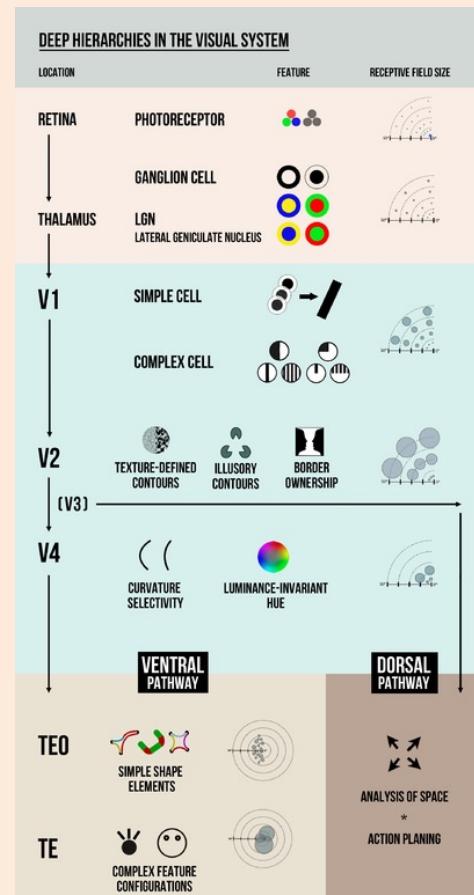
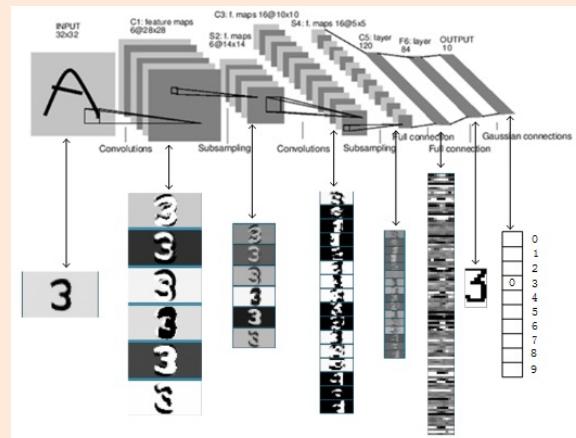
- 1950 and 1960s: Initial excitement.
 - Perceptron: linear classifier and stochastic gradient (roughly).
 - “the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.” New York Times (1958).
 - <https://www.youtube.com/watch?v=IEFRtz68m-8>
 - Object recognition assigned to students as a summer project
- Then drop in popularity:
 - Quickly realized **limitations of linear models.**



bonus!

ML and Deep Learning History

- 1970 and 1980s: **Connectionism** (brain-inspired ML)
 - Want “connected networks of simple units”.
 - Use **parallel computation** and **distributed representations**.
 - Adding hidden layers z_i increases expressive power.
 - With 1 layer and enough sigmoid units, a **universal approximator**.
 - Success in optical character recognition.



bonus!

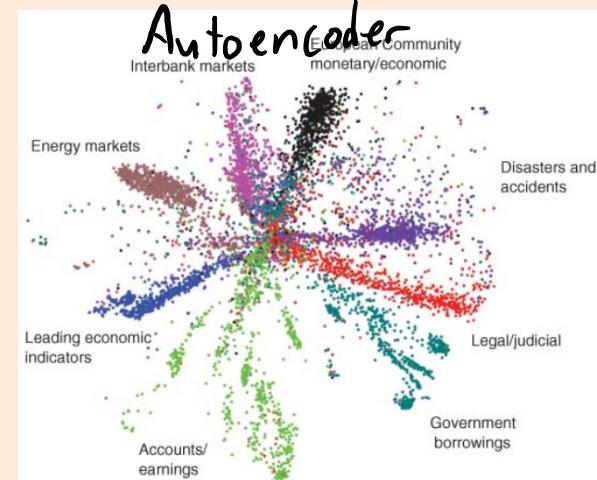
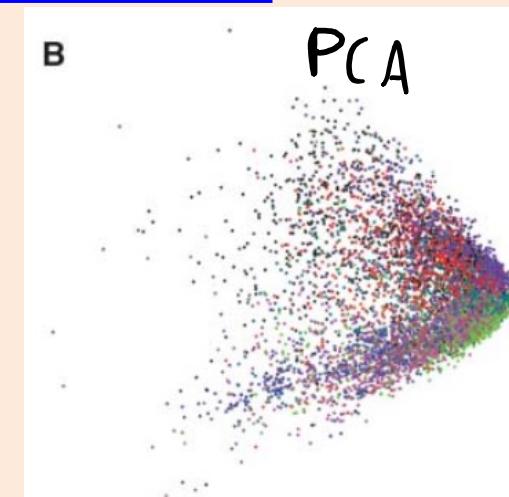
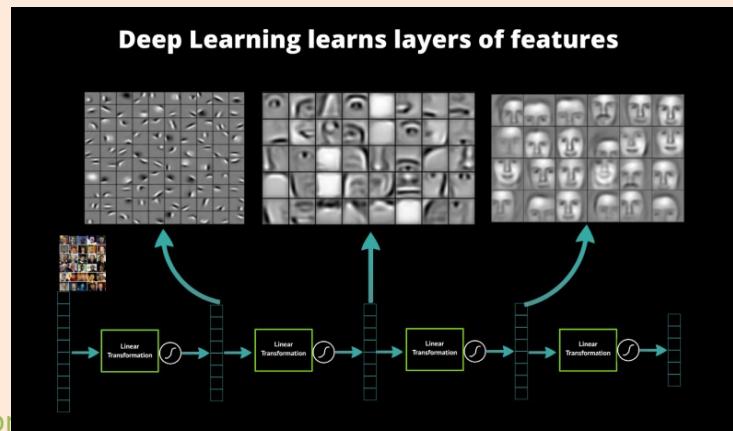
ML and Deep Learning History

- 1990s and early-2000s: drop in popularity.
 - It proved really difficult to get multi-layer models working robustly.
 - We obtained similar performance with simpler models:
 - Rise in popularity of logistic regression and SVMs with regularization and kernels.
 - Lots of internet successes (spam filtering, web search, recommendation).
 - ML moved closer to other fields like numerical optimization and statistics.

bonus!

ML and Deep Learning History

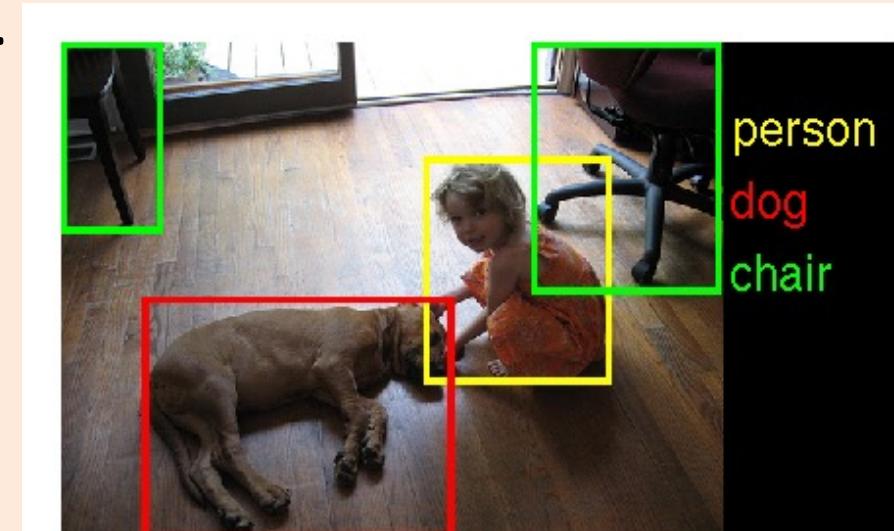
- Late 2000s: push to revive connectionism as “deep learning”.
 - Canadian Institute For Advanced Research (CIFAR) NCAP program:
 - “Neural Computation and Adaptive Perception”.
 - Led by Geoff Hinton, Yann LeCun, and Yoshua Bengio
 - Unsupervised successes: “deep belief networks” and “autoencoders”.
 - Could be used to initialize deep neural networks.
 - <https://www.youtube.com/watch?v=KuPai0ogiHk>



bonus!

2010s: DEEP LEARNING!!!

- Bigger datasets, bigger models, parallel computing (GPUs/clusters).
 - And some tweaks to the models from the 1980s.
- Huge improvements in automatic speech recognition (2009).
 - All phones now have deep learning.
- Huge improvements in computer vision (2012).
 - Changed computer vision field almost instantly.
 - This is now finding its way into products.



bonus!

2010s: DEEP LEARNING!!!

- Media hype:
 - “How many computers to identify a cat? 16,000”
New York Times (2012).
 - “Why Facebook is teaching its machines to think like humans”
Wired (2013).
 - “What is ‘deep learning’ and why should businesses care?”
Forbes (2013).
 - “Computer eyesight gets a lot more accurate”
New York Times (2014).
- 2015: huge improvement in language understanding.

Summary

- Neural networks learn features z_i for supervised learning.
- Sigmoid function avoids degeneracy by introducing non-linearity.
 - Universal approximator with large-enough ‘ k ’.
- Biological motivation for (deep) neural networks.
- Deep learning considers neural networks with many hidden layers.
 - Can more-efficiently represent some functions.
- Unprecedented performance on difficult pattern recognition tasks.
- Next time:
 - Training deep networks.

bonus!

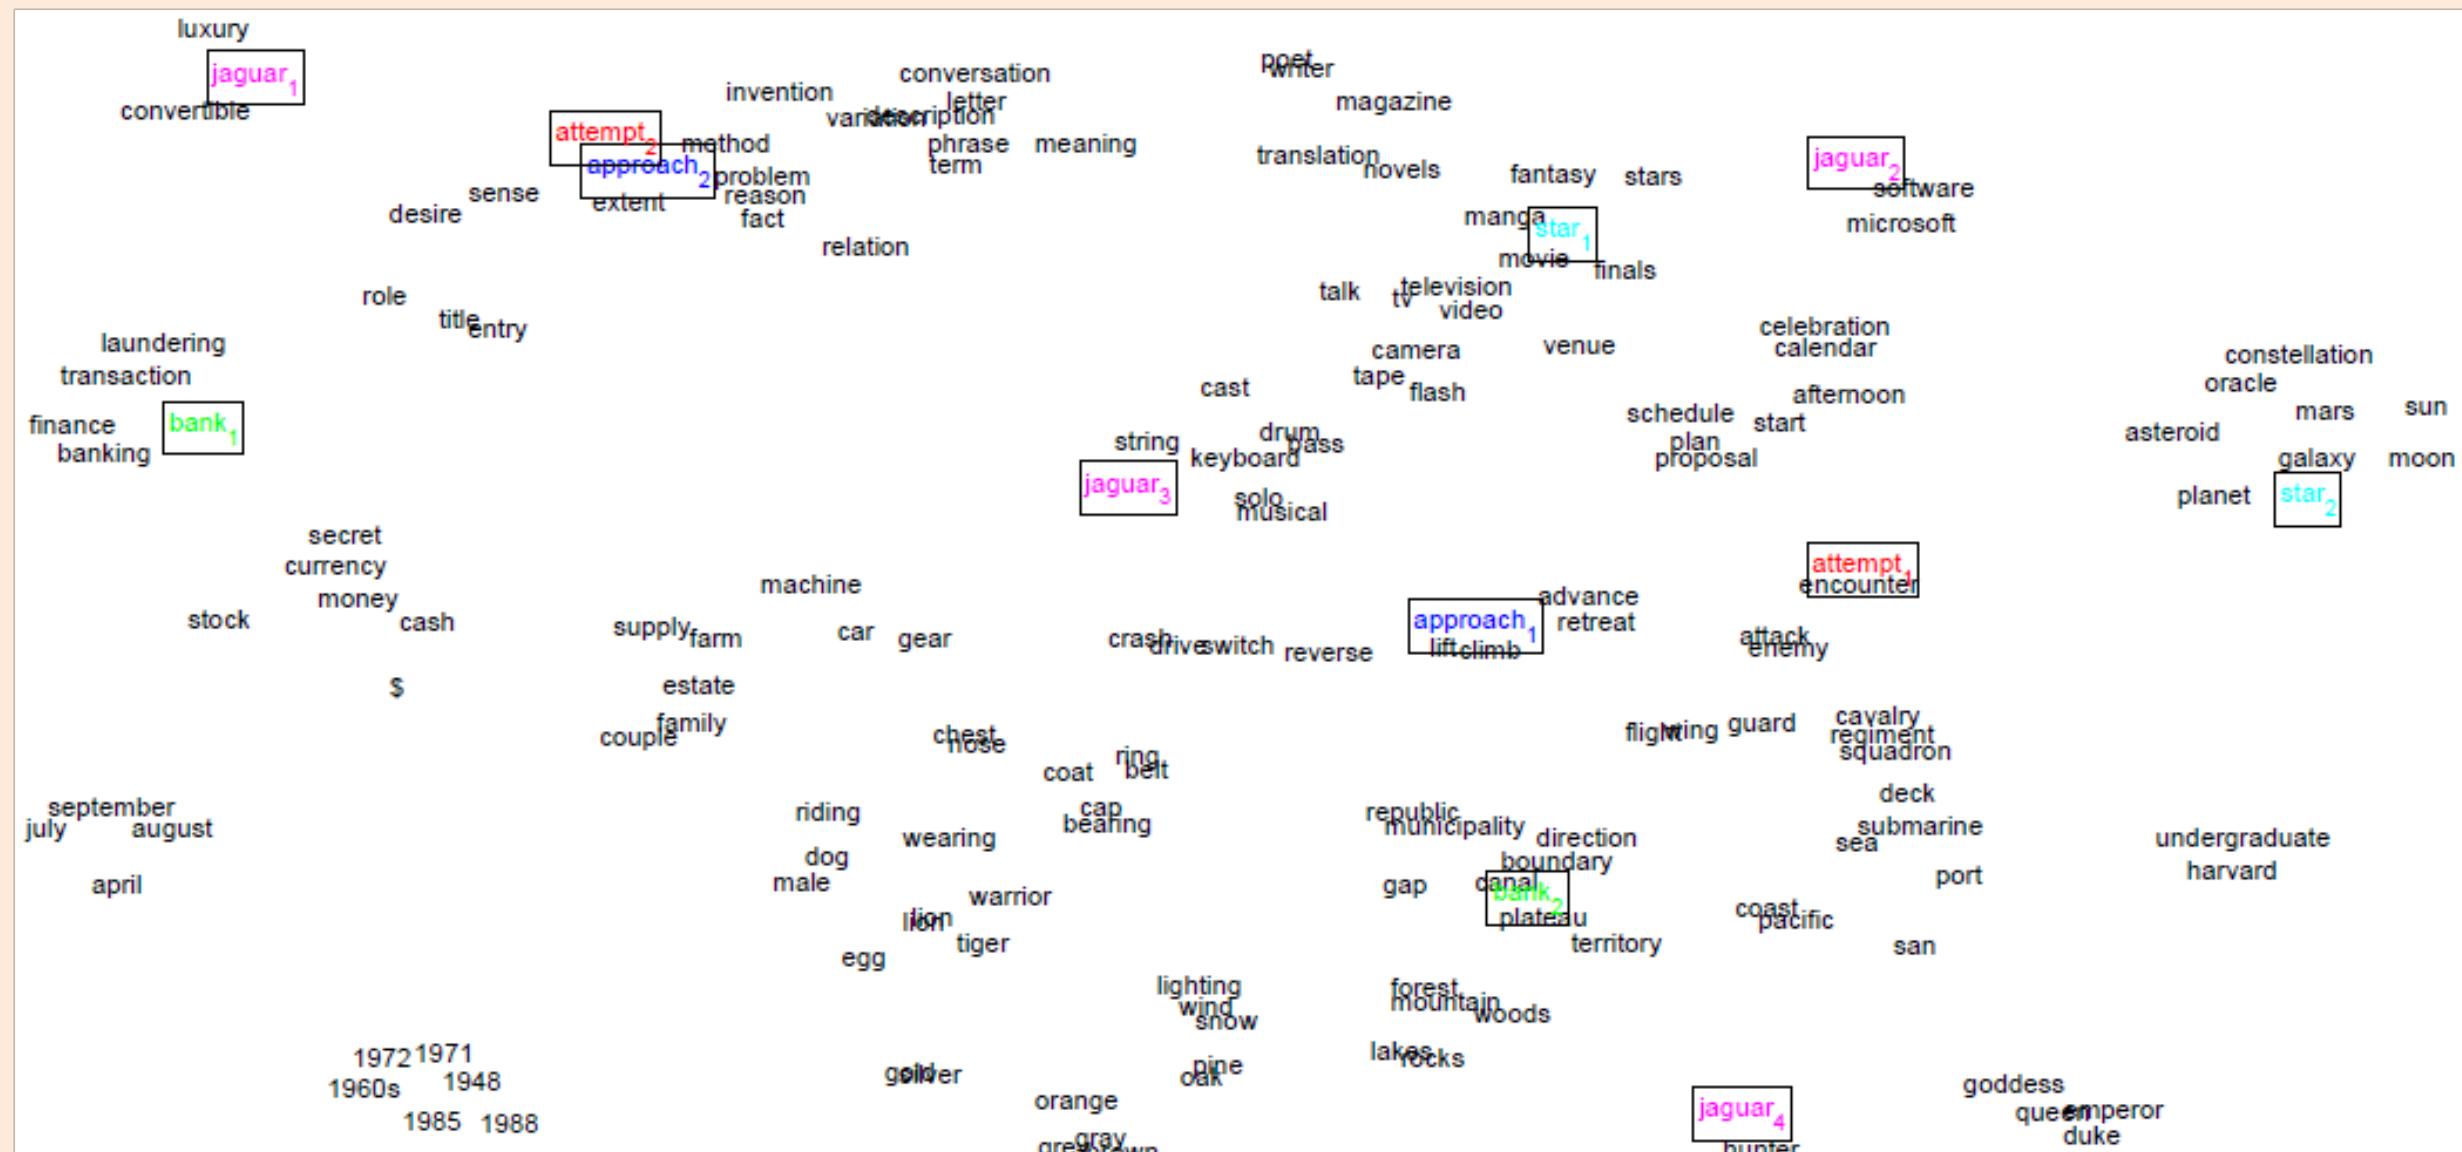
Multiple Word Prototypes

- What about **homonyms** and **polysemy**?
 - The word vectors would **need** to account for all meanings.
- More recent approaches:
 - Try to **cluster** the different contexts where words appear.
 - Use **different vectors** for different contexts.

$$X_{j\text{distr}} \approx \begin{bmatrix} \dots & \dots & \dots \\ \dots & \dots & \dots \\ \dots & \dots & \dots \end{bmatrix} \left. \right\} z_{j1} \\ \left. \right\} z_{j2} \\ \left. \right\} z_{j3}$$

bonus!

Multiple Word Prototypes



bonus!

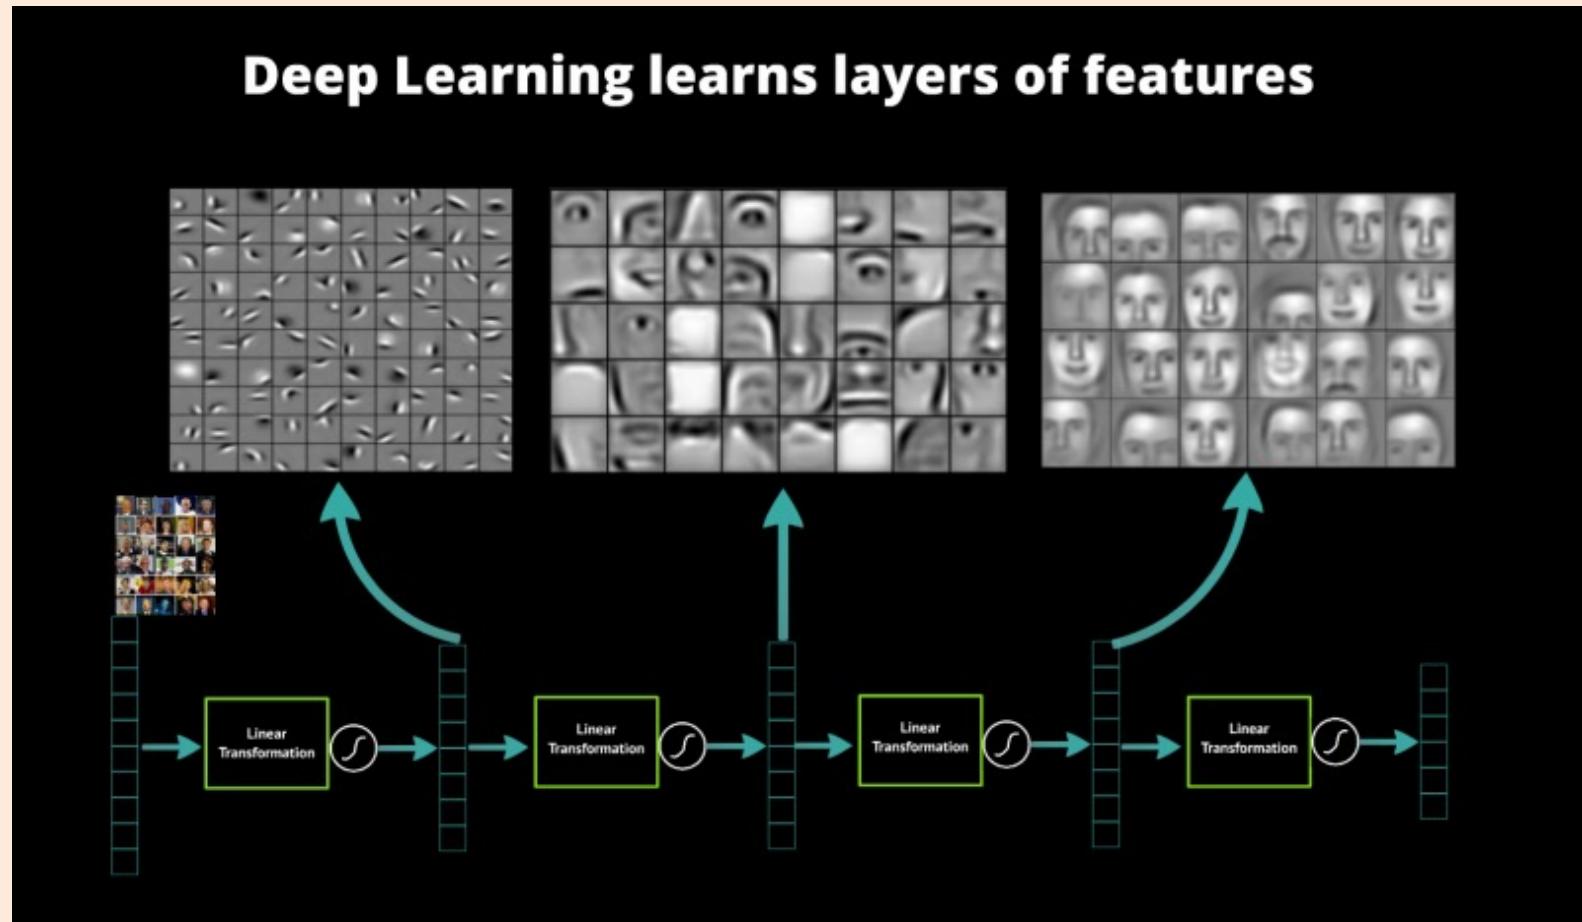
Why $z_i = Wx_i$?

- In PCA we had that the optimal $Z = XW^T(WW^T)^{-1}$.
- If W had normalized+orthogonal rows, $Z = XW^T$ (since $WW^T = I$).
 - So $z_i = Wx_i$ in this normalized+orthogonal case.
- Why we would use $z_i = Wx_i$ in neural networks?
 - We didn't enforce normalization or orthogonality.
- Well, the value $W^T(WW^T)^{-1}$ is just “some matrix”.
 - You can think of neural networks as just **directly learning this matrix**.

bonus!

Cool Picture Motivation for Deep Learning

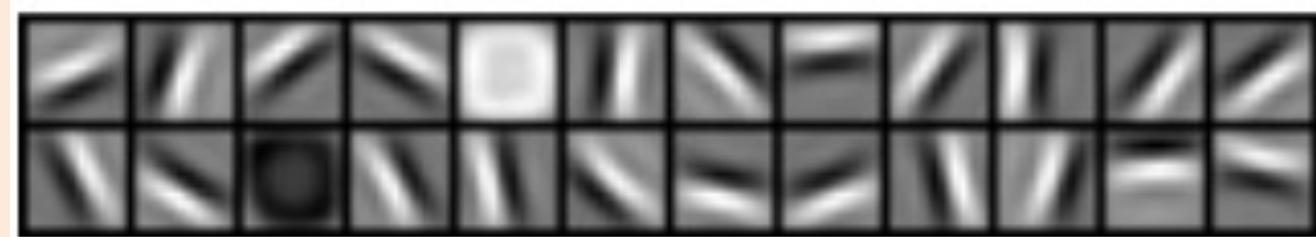
- Faces might be composed of different “parts”:



bonus!

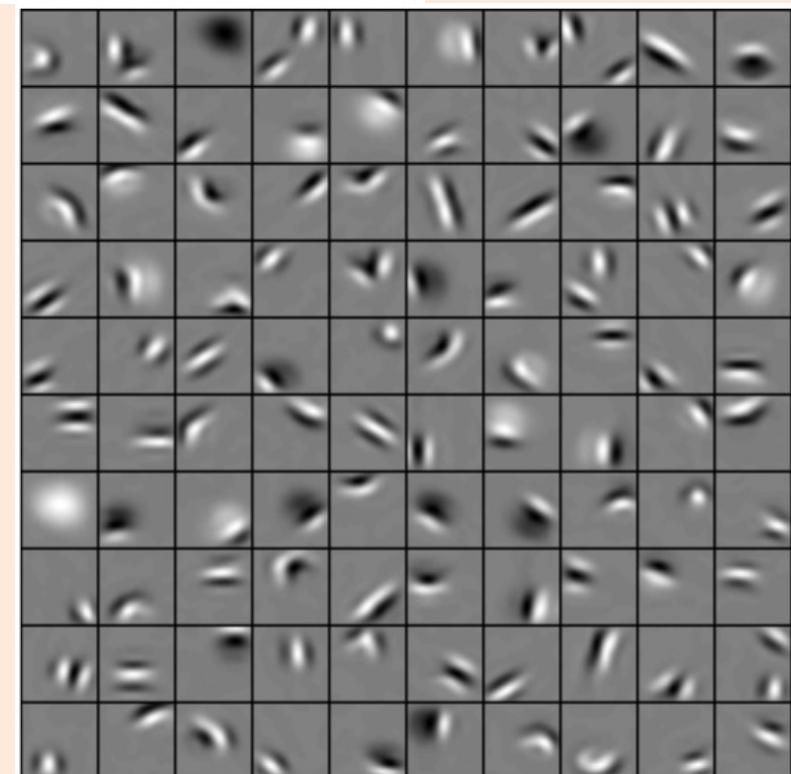
Cool Picture Motivation for Deep Learning

- First layer of z_i trained on 10 by 10 image patches:



} "Gabor filters"

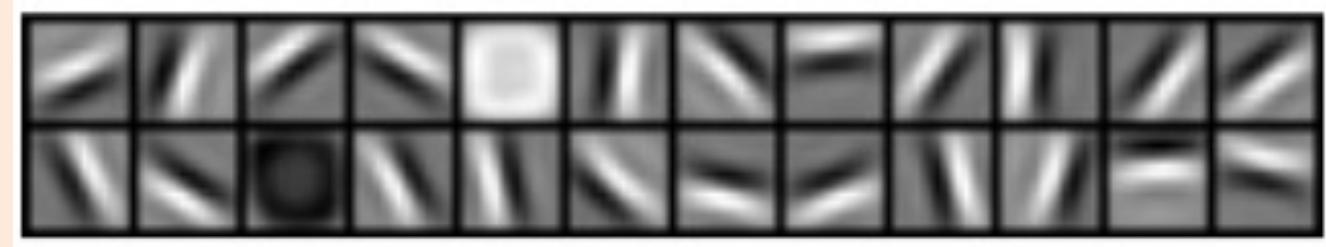
- Attempt to visualize second layer:
 - Corners, angles, surface boundaries?
- Models require many tricks to work.
 - We'll discuss these next time.



bonus!

Cool Picture Motivation for Deep Learning

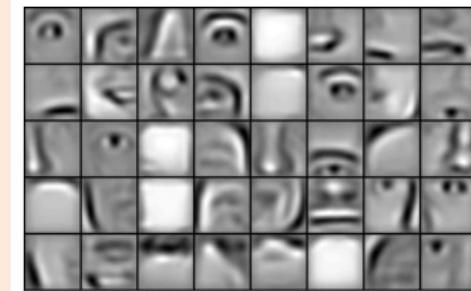
- First layer of z_i trained on 10 by 10 image patches:



} "Gabor filters"

- Visualization of second and third layers trained on specific objects:

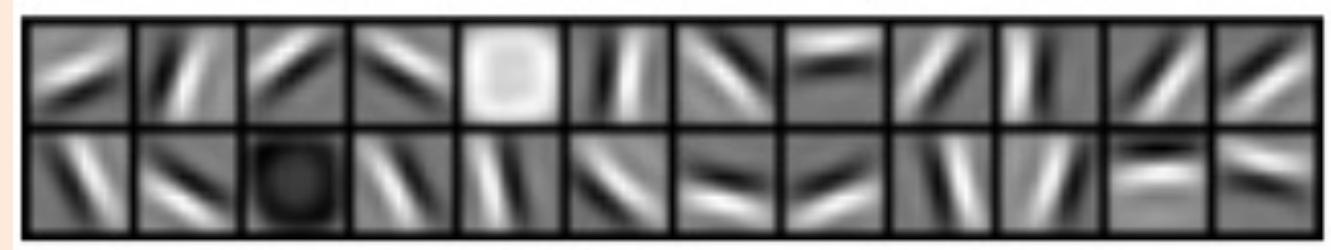
faces



bonus!

Cool Picture Motivation for Deep Learning

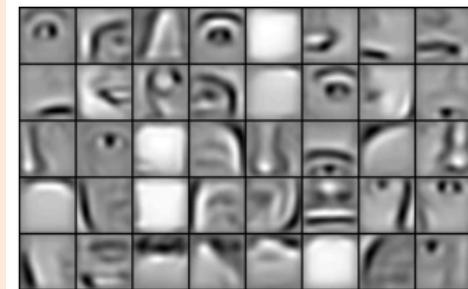
- First layer of z_i trained on 10 by 10 image patches:



} "Gabor filters"

- Visualization of second and third layers trained on specific objects:

faces



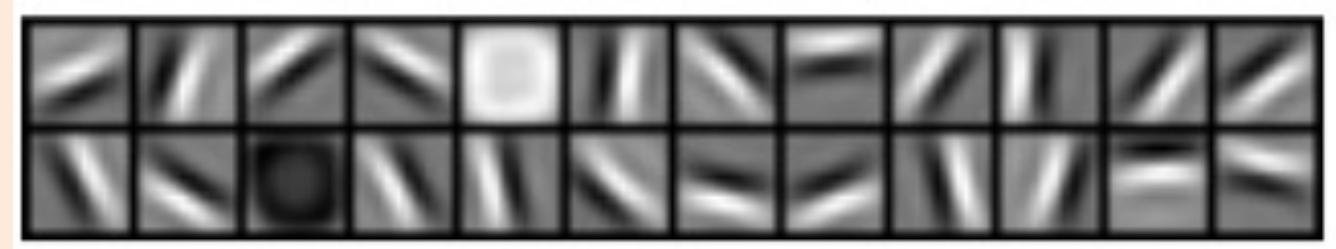
cars



bonus!

Cool Picture Motivation for Deep Learning

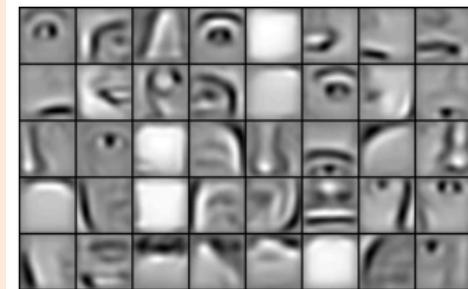
- First layer of z_i trained on 10 by 10 image patches:



} "Gabor filters"

- Visualization of second and third layers trained on specific objects:

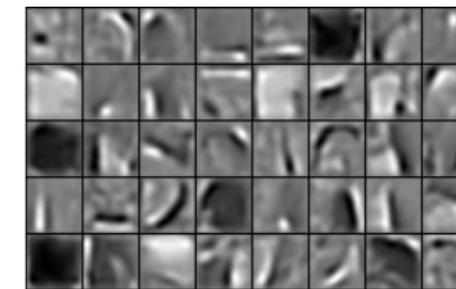
faces



cars



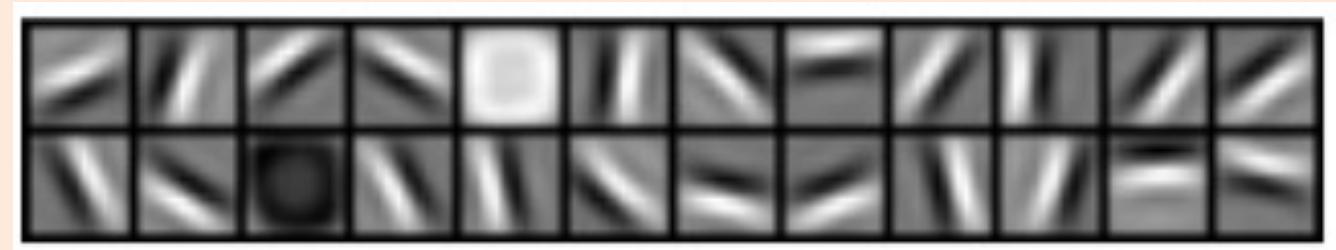
elephants



bonus!

Cool Picture Motivation for Deep Learning

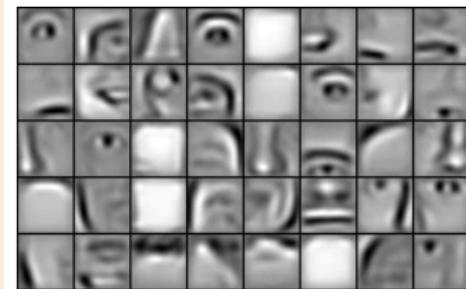
- First layer of z_i trained on 10 by 10 image patches:



} "Gabor filters"

- Visualization of second and third layers trained on specific objects:

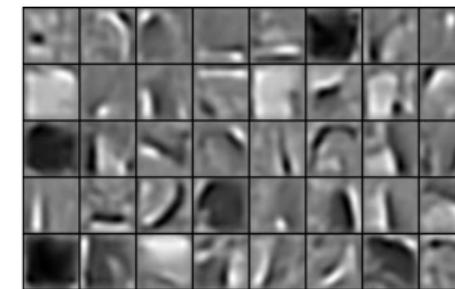
faces



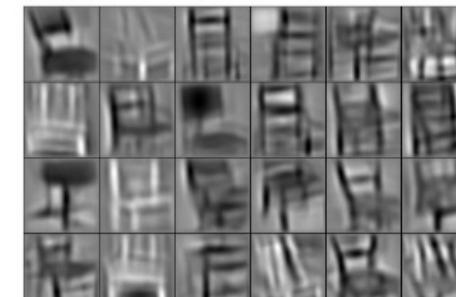
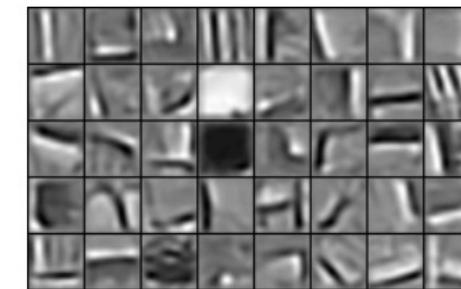
cars



elephants



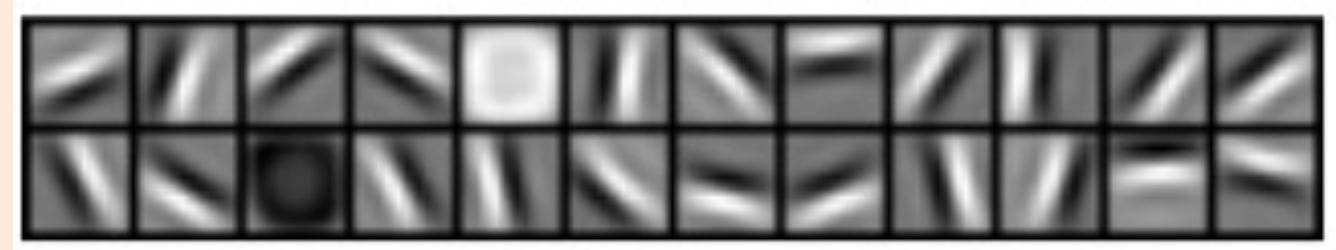
chairs



bonus!

Cool Picture Motivation for Deep Learning

- First layer of z_i trained on 10 by 10 image patches:



} "Gabor filters"

- Visualization of second and third layers trained on specific objects:

